

# HIGH PERFORMANCE Computing (HPC)



ALBERTO DASSATTI - 2023

WHAT IS GOING ON?

You need to measure!



**TO  
MEASURE  
IS TO  
KNOW**

A close-up, perspective view of a wooden ruler. The ruler is light brown with a visible wood grain. It features a scale with black markings and numbers from 1 to 11. The text 'TO MEASURE IS TO KNOW' is printed on the ruler in a bold, sans-serif font. 'TO' and 'IS TO' are in white, while 'MEASURE' and 'KNOW' are in bright yellow. The ruler is positioned diagonally across the frame, with the numbers 1 through 11 visible. The background is a plain, light gray surface.

**“In God we trust.  
All others must  
bring data”.**

W. Edwards Deming



A close-up photograph of a person's hands holding a metal vernier caliper. The caliper is held horizontally, with the beam facing the viewer. The word "Improvement" is printed in a large, bold, black sans-serif font across the center of the beam. The hands are positioned on either side of the caliper, with fingers gripping the frames. The background is a plain, light-colored surface.

**Improvement**

Every measure has its own tool. CS is not different.



And measuring correctly is even **harder**.

**Measuring time!**

# THE BASICS: TIME

bash

```
[al@lap ~]# time your_command  
[al@lap ~]# <your_command output omitted>  
[al@lap ~]# real    0m3.168s  
[al@lap ~]# user    0m2.952s  
[al@lap ~]# sys     0m0.180s
```

Let's test [hyperfine](#).



# THE BASICS: TIME

Can we have a little more visibility?

bash

```
[al@lap ~]# strace -tt
```

bash

```
[al@lap ~]# ltrace -tt
```

## THE BASICS: TIME

If you have access at the source code and the toolchain, some more options are available.

bash

```
[al@lap ~]# gcc -pg -O3 -Wall source.c -o source  
[al@lap ~]# ./source  
[al@lap ~]# gprof source gmon.out > analysis.txt
```

With the help of `-pg` there is also the option of [uftrace](#)

# MANUAL INSTRUMENTATION

Gcc can help.

rdtsc

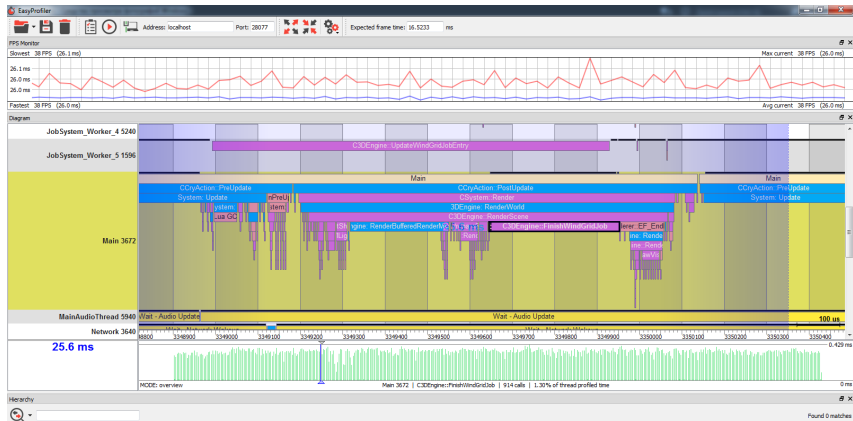
```
static __inline__ unsigned long long rdtsc(void)
{
    unsigned long long int x;
    __asm__ volatile (".byte 0x0f, 0x31" : "=A" (x));
    return x;
}
```

[source](#)

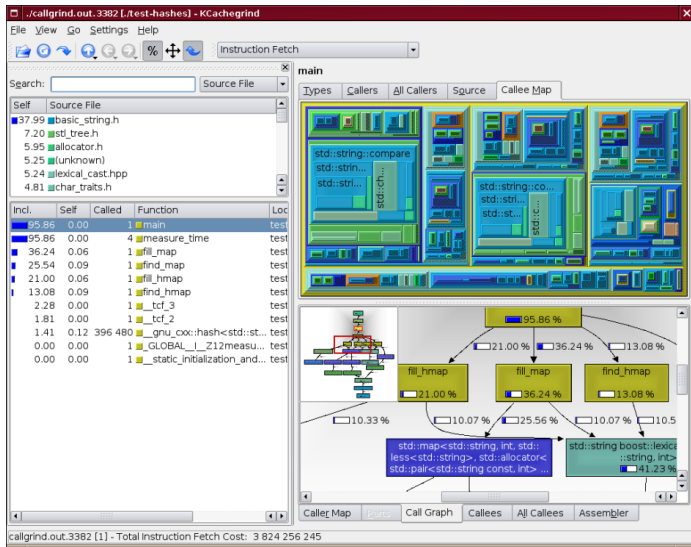
[cpucycles](#) is a library that improve a lot on this raw technique.

## MANUAL INSTRUMENTATION ON STEROIDS

You are not forced to trace all the code, but if you know your critical path you can use an instrumentation library (eg. [easy\\_profiler](#) or [Likwid](#) marker API).



# RUN IT ON AN EMULATOR: CALLGRIND, CACHEGRIND



a good tutorial

## WARNING AND LIES

Understand the numbers provided by the tools is not always easy. Please read this [article](#) and this [discussion](#) to be sure you get the most accurate info from your tool. I strongly recommend you this fun and instructive [video](#).

---

**What our system is doing?**

# WHAT OUR SYSTEM IS DOING

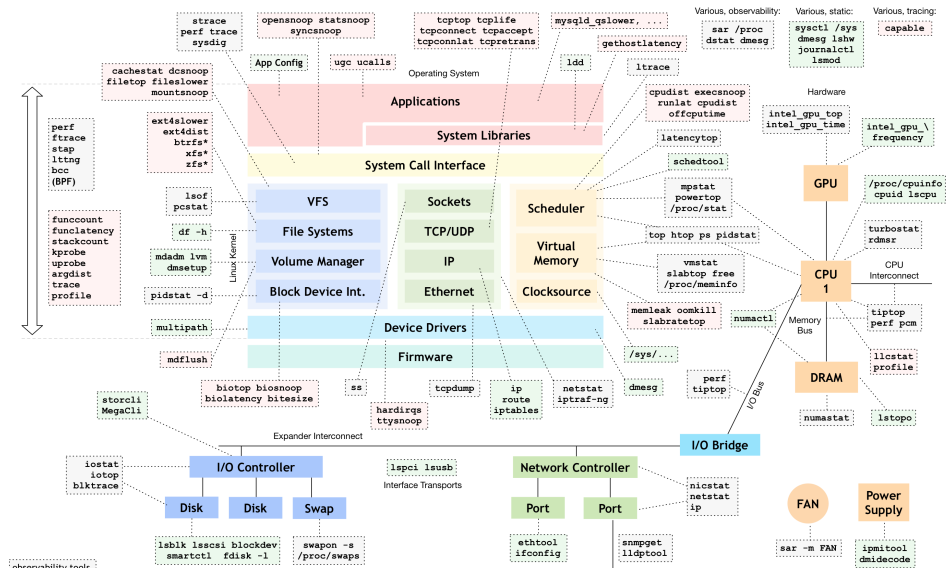
For a first look we may use

```
[al@lap ~]# top
```

bash

htop is a feature packed alternative





these can observe the state of the system at rest, without load

<https://github.com/brendangregg/perf-tools>
<https://github.com/iovisor/bcc>

style inspired by reddit.com/u/redct

<http://www.brendangregg.com/linuxperf.html>
2017

Up to date tools catalog can be found [here](#) and [here](#).

OK, BUT WHY IT IS TAKING SO LONG?

# Workload Characterization!

### micro Benchmarking

Isolate a small kernel or a function and analyze it alone

vs

### System Benchmarking

Looking at the performance of a system fully in service

## Statistical

Periodically stop the system and look at its state. If you collect enough data you can have a good picture of the situation.

- ▶ Advantage: do not need any code modification
- ▶ Warning: can be misleading, use long runs
- ▶ Example: `perf`

vs

## Instrumentation

Modify your code manually or automatically to emit `events` and then process them

- ▶ Advantage: tracing can be very precise
- ▶ Warning: you cannot see what's outside of your code or what your input is not stimulating
- ▶ Example: `gprof`

## Time-based sampling (TBS)

Program a timer. When the timer ends look at the `pc` and at the *stack* to get useful information.

- ▶ It's useful to have the debug symbols (-g)
- ▶ Trade sampling frequency for accuracy (limit the overhead)
- ▶ Some more ideas (-fno-exceptions, -fno-rtti, -fno-omit-frame-pointer)

vs

## Event-based sampling (EBS)

Count some events (HW or SW generated) and try to get what's going on.

- ▶ Low overhead (Intel < 2% using performance monitor counters)
- ▶ Use event ratios (cache miss/ cache access)

# PERFORMANCE COUNTERS

## oprofile, perf, VTune

- ▶ Total instruction count and mix
- ▶ Branch events
- ▶ Load/store events
- ▶ L1/L2 cache events
- ▶ Prefetching events
- ▶ TLB events
- ▶ Multicore events

Use **event ratios**

```
sudo perf list
```

# Perf Tutorial

[blog](#), [blog](#), [pdf](#), [Cern](#)



# PERF GOODIES

- ▶ Pmc-cloud-tools
- ▶ HotSpot
- ▶ Flamegraph

# PERF IS AWSOME

but, what if what you are interested in is not already an event or you want more visibility?

Dynamic instrumentation is the answer!

## UPROBE AND KPROBE

Since long time (3.15 at least), Linux has the infrastructure to dynamically add trace point in both kernel space (kprobes) and user space (uprobe).

They are not usually used directly. `perf probe` can help.

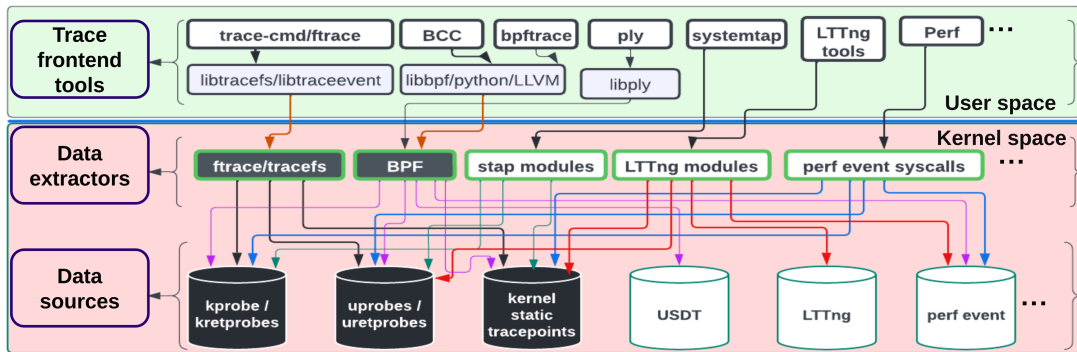
# Tracing

Bootlin Debugging-slides from pg 144

# BEHIND THE SCENE

## Linux trace system

4



From [Here](#).

# TRACING



ftrace



perf\_events



eBPF



SystemTap



LTTng



ktap



dtrace4linux



OEL DTrace



sysdig

To know more:

- ▶ have a look [here](#)
- ▶ read carefully and understand these [Methods](#)
- ▶ dig [this](#) mine of information

# QUESTIONS

