

HIGH PERFORMANCE Co<sup>mputing</sup><sub>ding</sub> (HPC)

ALBERTO DASSATTI - 2024



# WHO

- ▶ Professor: Alberto Dassatti
- ▶ Office: A11 (Cheseaux)
- ▶ E-mail: [alberto.dassatti@heig-vd.ch](mailto:alberto.dassatti@heig-vd.ch)
- ▶ Telephone: +41 24 557 61 60

2002: Master in Electronics Engineering, Politecnico di Torino

2003: Interim at Advanced System Technologies  
STMicroelectronics, Geneva

2004-2008 : Ph.D. in Telecommunications and Electronics  
(reconfigurable systems), Politecnico di Torino,

2008 : Visiting Ph.D. student UNSW Sydney, Australia

2008-2009: Post-Doc Researcher at VLSI Lab, Politecnico di Torino


2003-2010: Entrepreneur fondateur of MicroC s.n.c., a  
Consultancy firm

2010-2012: ingénieur at NATO Undersea Research Centre


2012-2013: ingénieur au REDS

2013-: Professeur HES HEIG-VD

2018-: REDS' Director



# WHAT YOU KNOW



WHAT IS HPC GOAL?

Make programs run faster!

# WHAT IS HPC ABOUT

1. Defining *performances*

# WHAT IS HPC ABOUT

1. Defining *performances*
2. Measure *performances*

# WHAT IS HPC ABOUT

1. Defining *performances*
2. Measure *performances*
3. Some modern Hardware (Super Scalar CPU, Multi Cores, Caches, SIMD, GPU)



# WHAT IS HPC ABOUT

1. Defining *performances*
2. Measure *performances*
3. Some modern Hardware (Super Scalar CPU, Multi Cores, Caches, SIMD, GPU)
4. Something about Compilers (LLVM, gcc)

# WHAT IS HPC ABOUT

1. Defining *performances*
2. Measure *performances*
3. Some modern Hardware (Super Scalar CPU, Multi Cores, Caches, SIMD, GPU)
4. Something about Compilers (LLVM, gcc)
5. Something about libraries and languages (CUDA, OpenCL, BLAS)

# WHAT IS HPC ABOUT

1. Defining *performances*
2. Measure *performances*
3. Some modern Hardware (Super Scalar CPU, Multi Cores, Caches, SIMD, GPU)
4. Something about Compilers (LLVM, gcc)
5. Something about libraries and languages (CUDA, OpenCL, BLAS)
6. Become a better programmer


WHY

NEWS

## Data centers are the new polluters



report



Source: Intel Corp.

# WHY



# DOES ANYONE REALLY CARES?

CppCon 2014: Andrei Alexandrescu

# Performance Summary


## The BIG Picture

$$\text{CPU Time} = \frac{\text{Instructions}}{\text{Program}} \times \frac{\text{Clock cycles}}{\text{Instruction}} \times \frac{\text{Seconds}}{\text{Clock cycle}}$$

- Performance depends on
  - Algorithm: affects IC, possibly CPI
  - Programming language: affects IC, CPI
  - Compiler: affects IC, CPI
  - Instruction set architecture: affects IC, CPI,  $T_c$




# DEEPER REASONING



# End of Growth of Performance?

## 40 years of Processor Performance



Based on SPECintCPU. Source: John Hennessy and David Patterson, Computer Architecture: A Quantitative Approach, 6/e. 2018

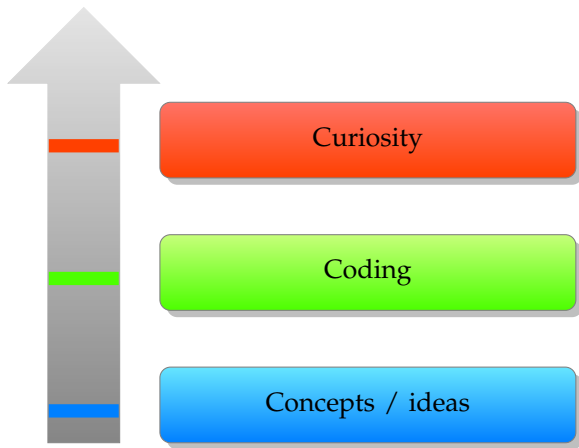
How



# How

- ▶ Theory presentations: 2h/w
- ▶ Lab work: 2h/w
- ▶ Your research and ideas
- ▶ May be some expert seminars
- ▶ Your personal effort (reading and watch video)
- ▶ Bruno and me are here to support and help (ask a meeting by email)
- ▶ It's all up to you

# How

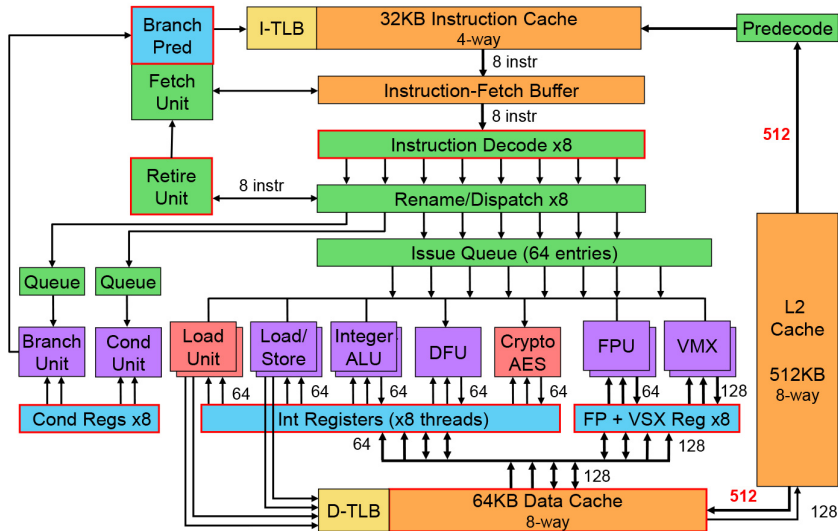


# CREATE A MENTAL MODEL

- ▶ Before looking at code
- ▶ Model the principal components
- ▶ Set clear expectations:
  - ▶ Back of the envelop calculations  
<https://github.com/sirupsen/napkin-math>
  - ▶ Look for similar systems
- ▶ Refine your model with all available information (bayesian reasoning)
- ▶ Don't hate statistics...

# FROM C ABSTRACT MACHINE TO HARDWARE

`i += 3;`



# WHAT YOU NEED

## to know

- ▶ System Architecture (the basics)
- ▶ Operating System (SYE)++
- ▶ Programming in C (very well)
- ▶ A little of assembly is useful
- ▶ Linux
- ▶ version control (git, ...)

## and...

- ▶ Curiosity
- ▶ Adapt your coding style
- ▶ Investigate and relentless improving
- ▶ read, read, test, read, ask, test, read again
- ▶ a little statistics ...




# IT IS HARD

- ▶ Very complex situations (each system is different)
- ▶ Hard to work in isolation, reproduce
- ▶ Custom tools
- ▶ Concurrency (SW and HW/SW)
- ▶ Few methodologies

ET SURPRISING AT TIMES

Unusual Disk Latency



## Why C?


- ▶ Portable
  - ▶ You are the master
  - ▶ A lot of tools
  - ▶ A lot of libraries
  - ▶ Other languages can call high performance C libs
- and
- ▶ OSes are written in C and a little assembler (yes, Windows and OSX too)
  - ▶ [GitHub](#)
  - ▶ [Popularity](#)


# C CAN SURPRISE YOU IN MANY WAYS

- ▶ It's very **influential**
- ▶ It's **complex**
- ▶ if you look at performances can **shock** you
- ▶ Cay Horstmann **presentation** is instructive

## Why Linux?

- ▶ It's documented
- ▶ It's open
- ▶ Tons of tools
- ▶ Tons of libraries
- ▶ It's very **common**





## Please, please and please

Write readable, litterate, clear  
and meaningful code!

Looking for performance cannot  
be used as an excuse for bad  
written code.


Some software engineering Best  
Practices can be derogated.

# WHERE WE OPTIMIZE



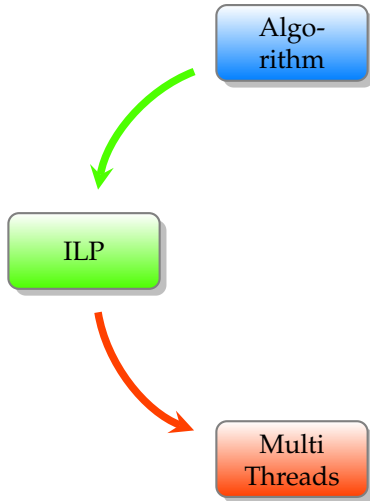
Algo-  
rithm

# WHERE WE OPTIMIZE







# WHERE WE OPTIMIZE



## WHERE WE OPTIMIZE



## WHERE WE OPTIMIZE



## REFERENCES

<https://datacenters.lbl.gov/resources/dc-power-improved-data-center-efficiency>

# QUESTIONS

