



HE^{VD}
IG

Intelligence Artificielle pour les systèmes autonomes (IAA)

End-to-end learning and direct perception

Prof. Yann Thoma - Prof. Marina Zapater

Février 2024

Basé sur le cours du Prof. A. Geiger



Outline

Today's lesson

→ Imitation Learning

- PilotNet and Dronet
- Conditional Imitation Learning

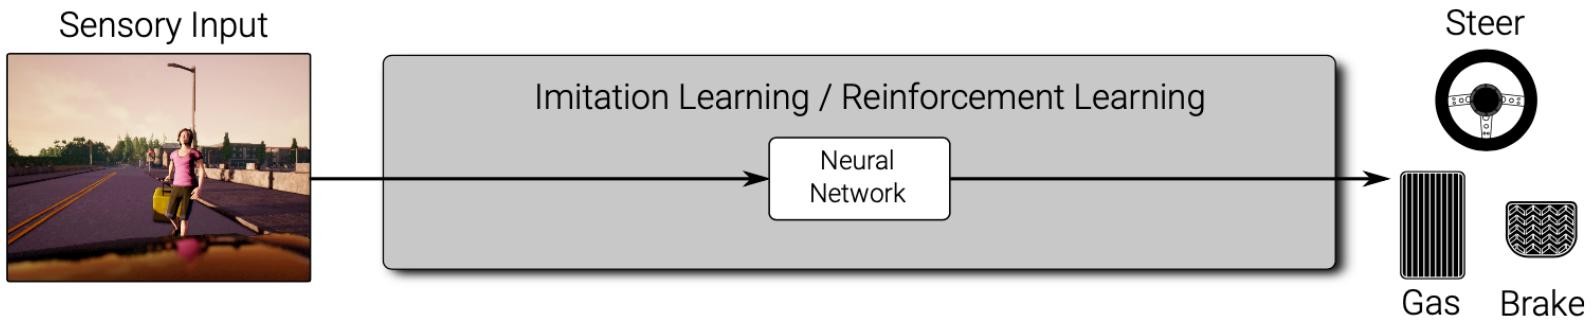
→ Direct perception

- Affordance Learning
- Conditional Affordance Learning
- Visual abstractions



End-to-end learning

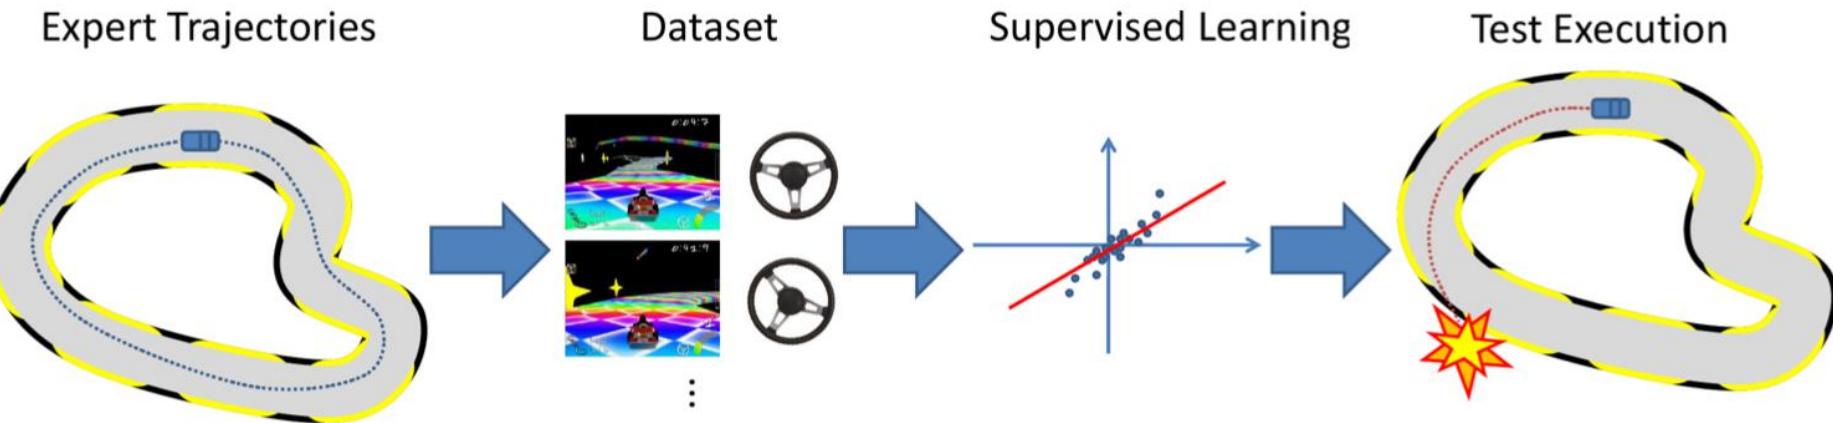
The “Second approach” we saw on our introduction lecture



- End-to-end requires cheap(er) annotations
- Issues with training/generalization
- Lack of interpretability...
- Certification issue...

Imitation learning basics

Learning from “experts”



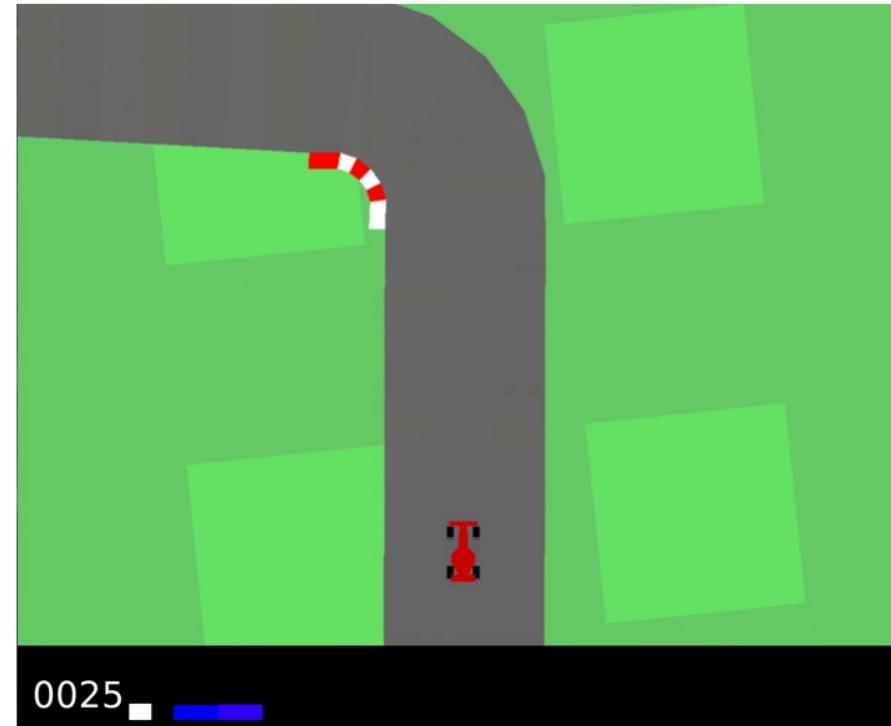
- Hard-coding policies is difficult : better use a data-driven approach
- Given demonstrations (expert trajectories) train a policy to mimic decisions

Imitation Learning

Car Racing: Imitating the human driver's trajectories



Trainer
(Human Driver)

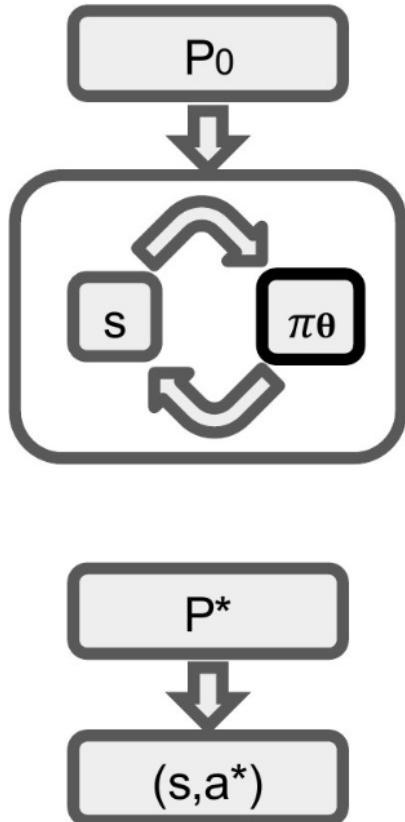


Trainee
(Neural Network)

A “more formal” definition of imitation learning

For general imitation learning

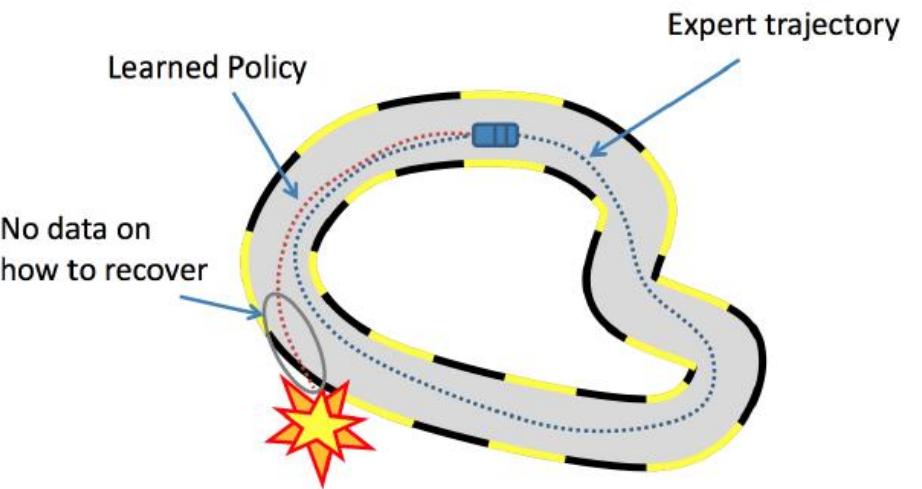
- We have an environment (Markov Decision process) comprised of:
 - States (s) : what can be observed (for example the game screen)
 - Actions (a) : the actuations we can perform (e.g.: speed, turning angle)
- The goal is to learn a policy π_θ
 - The complexity relies on the fact that each time the expert takes an action, that brings you to a different state...
 - So you would need to ask the expert each and every time, or said differently the state distribution $P(s|\pi_\theta)$ depends on the roll-out of the policy
- To solve this problem we use “behavioral cloning”
 - The state distribution is reduced P^* to the possibilities seen by the expert
 - Now this is a supervised learning problem



The simplest “flavour” of imitation learning

Behavioural cloning (when it works and when it does not)

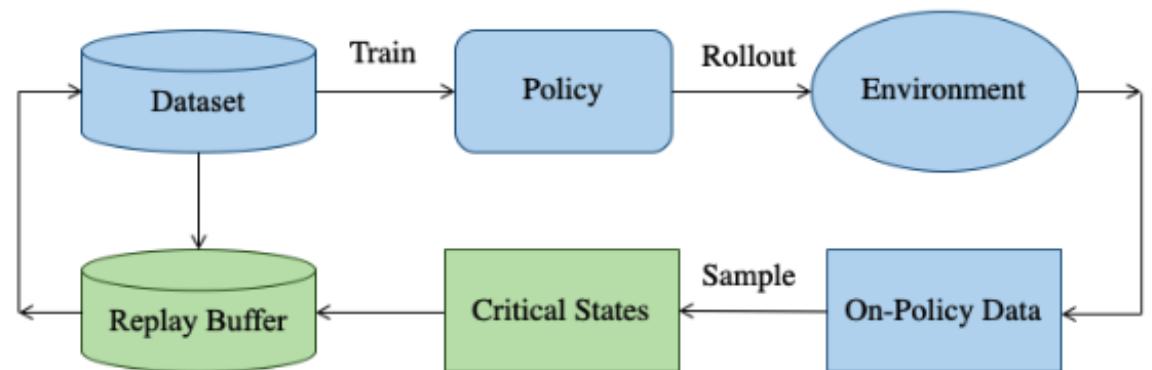
- We try to find a policy that given a certain state provides the best action $(s_0, a_0), (s_1, a_1)...$
- At each decision interval, we minimize the loss function $L(a^*, \pi_\theta(s))$
- What if π_θ makes a mistake?
 - We'll enter into a state we never saw before
 - We have no data from the expert on how to recover
 - We cannot recover (we crash!)
- So... what do we do?



Training with critical states

Having access to an interactive demonstrator: data and policy aggregation

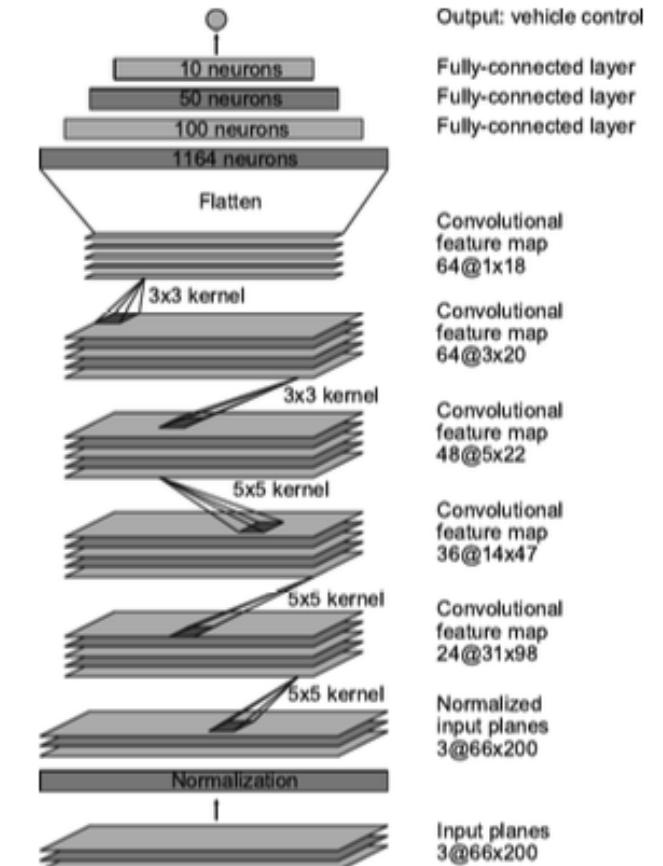
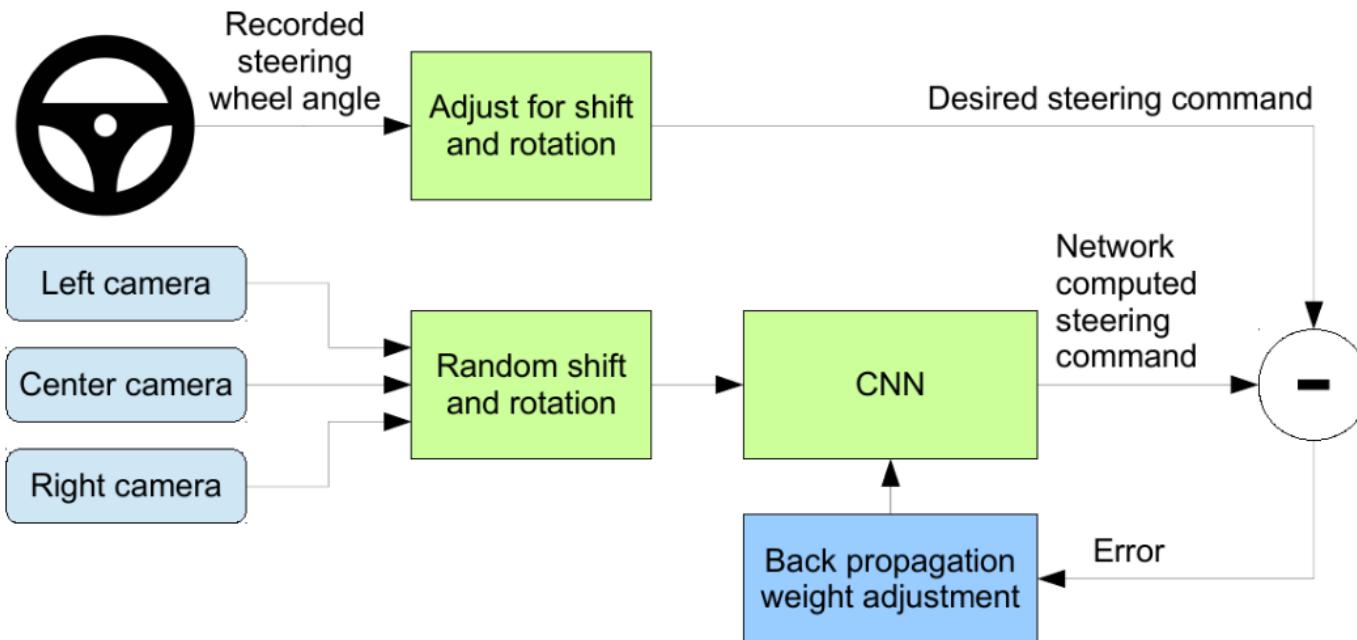
- Sample critical states from the collected policy data based on the utility they provided to the learned policy (in terms of driving behavior)
- Incorporate a replay buffer which focuses on the high uncertainty regions
- Or... you choose another algorithm



PilotNet: End-to-end learning for self-driving cars

System overview

- 3 cameras provide virtually shifted/rotated images
- Convolutional network (250k params)
- YUV image representation



PilotNet

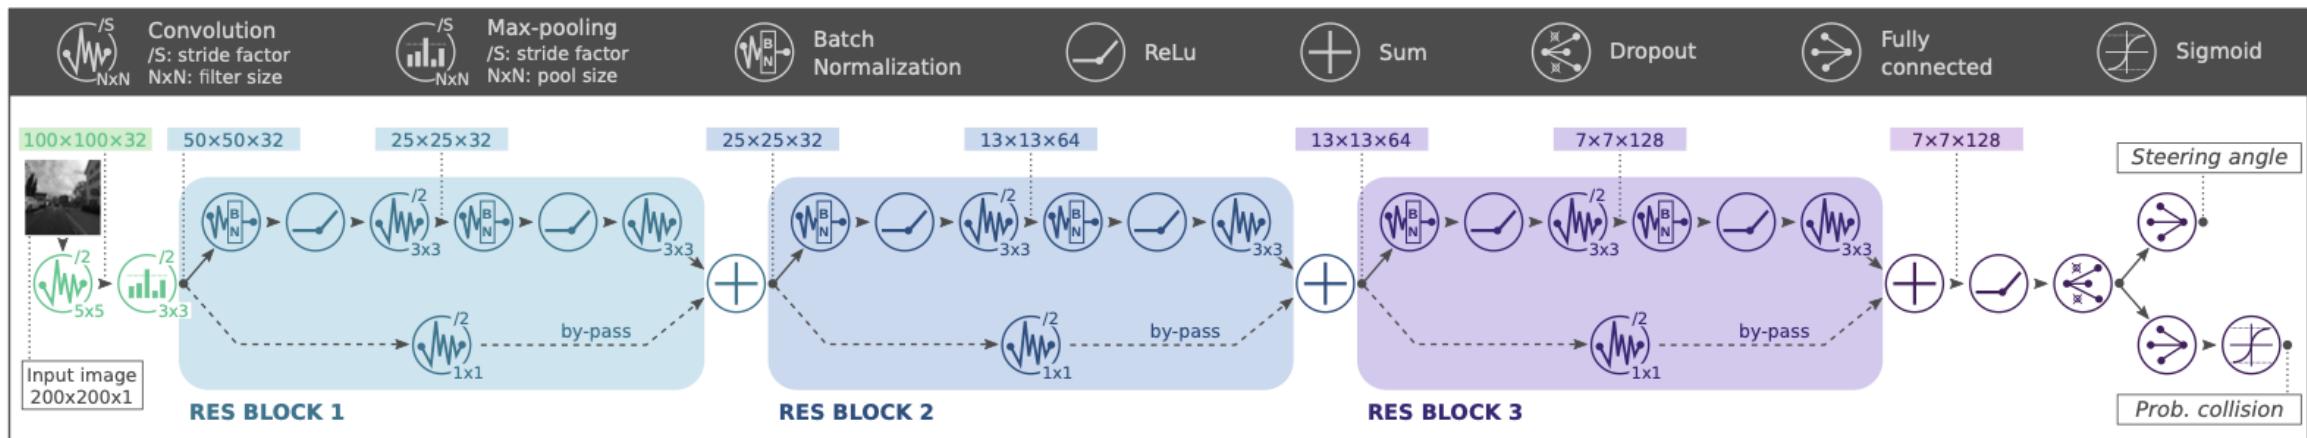
NVIDIA BB8



Dronet

And simplifications over Dronet to make Crazyflie fly

- Using input images from the drone camera (grayscale)
- Feeding those to ResNets
- To obtain the probability of collision and the steering angle the drone should take



Dronet Training and Testing

And performance

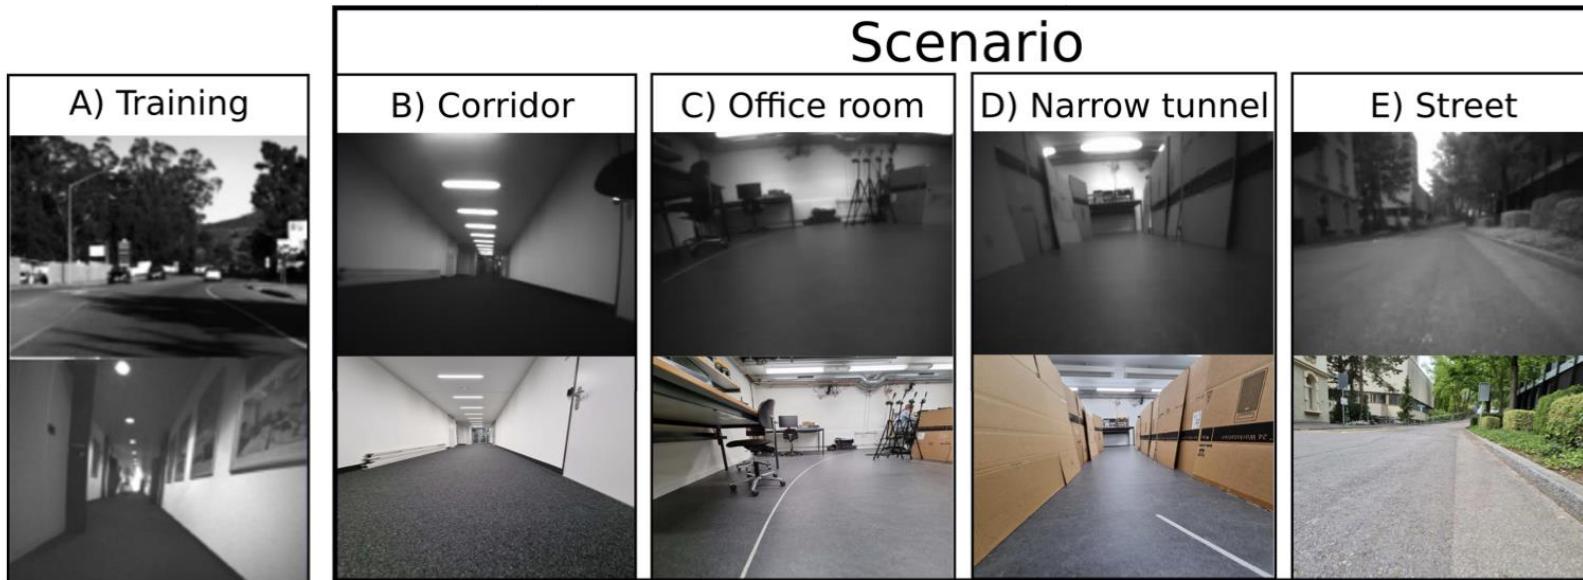


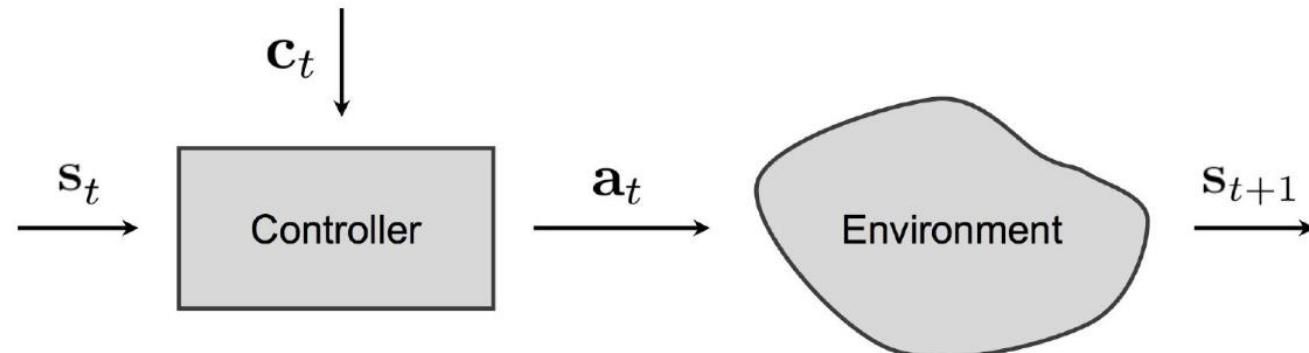
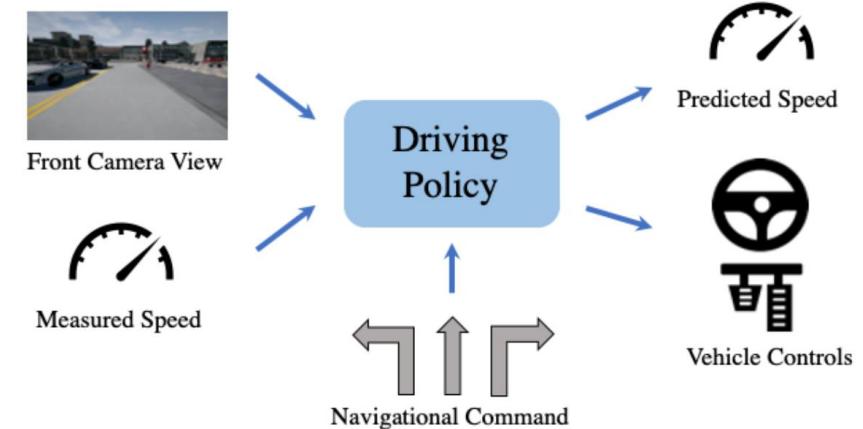
Fig. 8. Samples of: A) the training images (both Udacity and Himax dataset); B) working-place corridor; C) office room with furniture; D) narrow pipeline-like tunnel; E) public street.



Conditional imitation learning

Taking decisions when there are several possibilities

- We add a condition controller on the navigation command (left, right, straight)
- High-level navigation command provided by GPS
 - Telling the vehicle where to go on next turn (left, right, straight)
 - Reducing ambiguity induced by the environment



Conditional imitation learning

End-to-end driving via Conditional Imitation Learning



(a) Aerial view of test environment

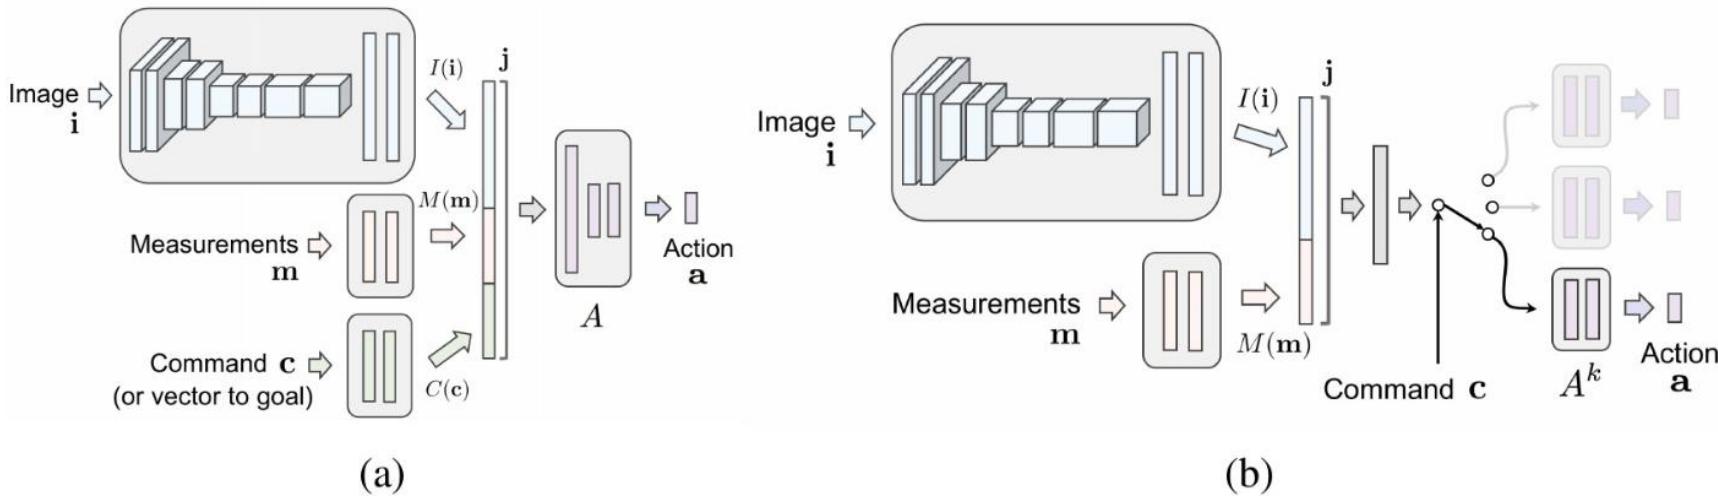
(b) Vision-based driving, view from onboard camera

(c) Side view of vehicle

<https://www.youtube.com/watch?v=cFtnflNe5fM>

Different proposals for conditional imitation learning

Two example architectures

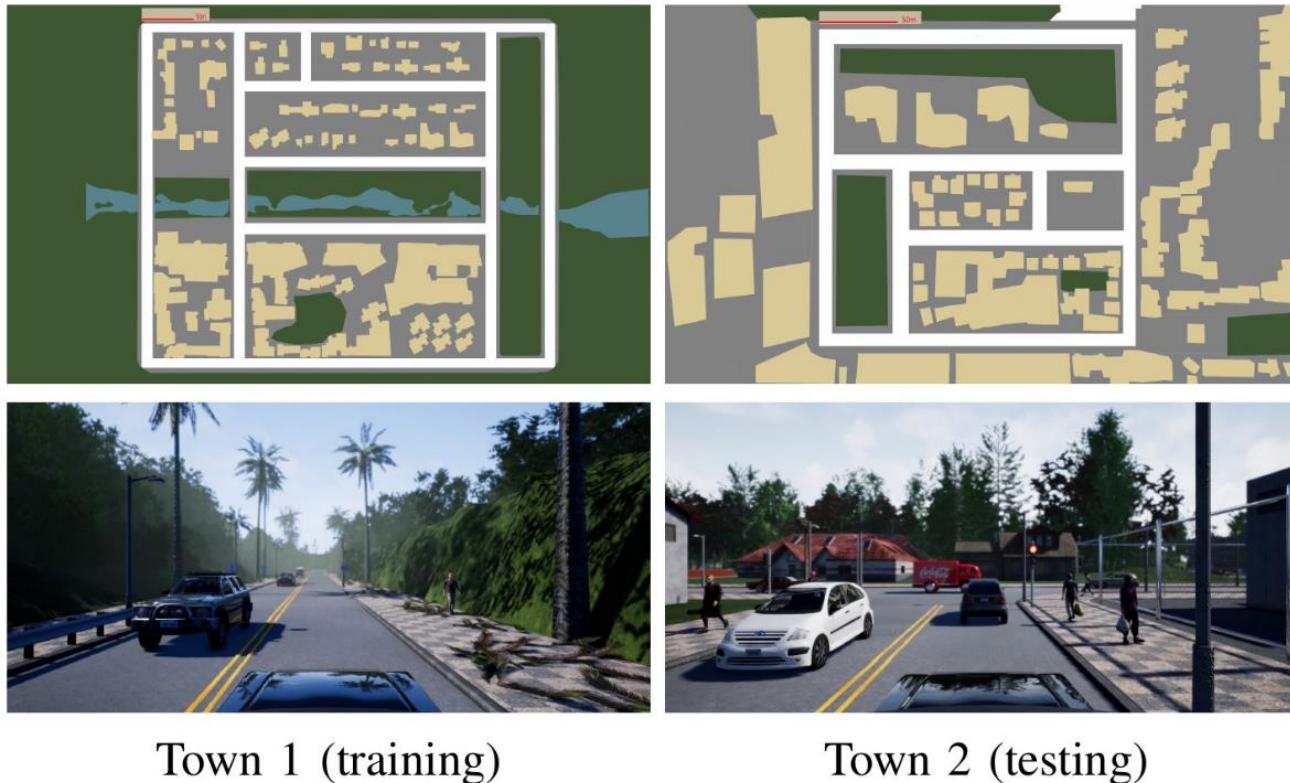


Measurements “m” capture additional information (speed of vehicle)

- a) Extract features and concatenate with image features $C(c)$
- b) Command “c” acting as a switch between specialized submodules

Code available for testing on Carla simulator

Training in one town, testing on another



Town 1 (training)

Town 2 (testing)

Outline

Today's lesson

→ Imitation Learning

- PilotNet and Dronet
- Conditional Imitation Learning

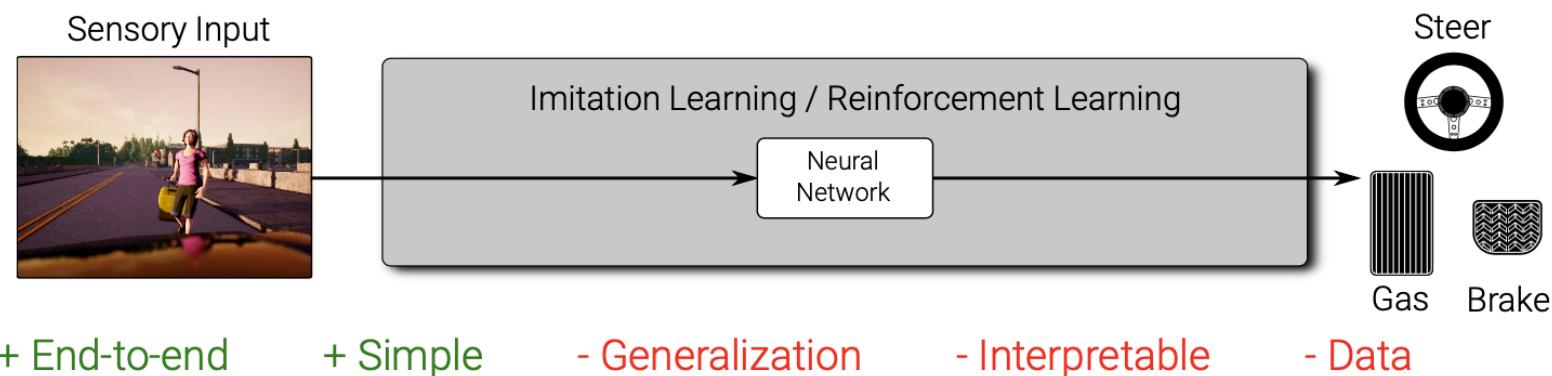
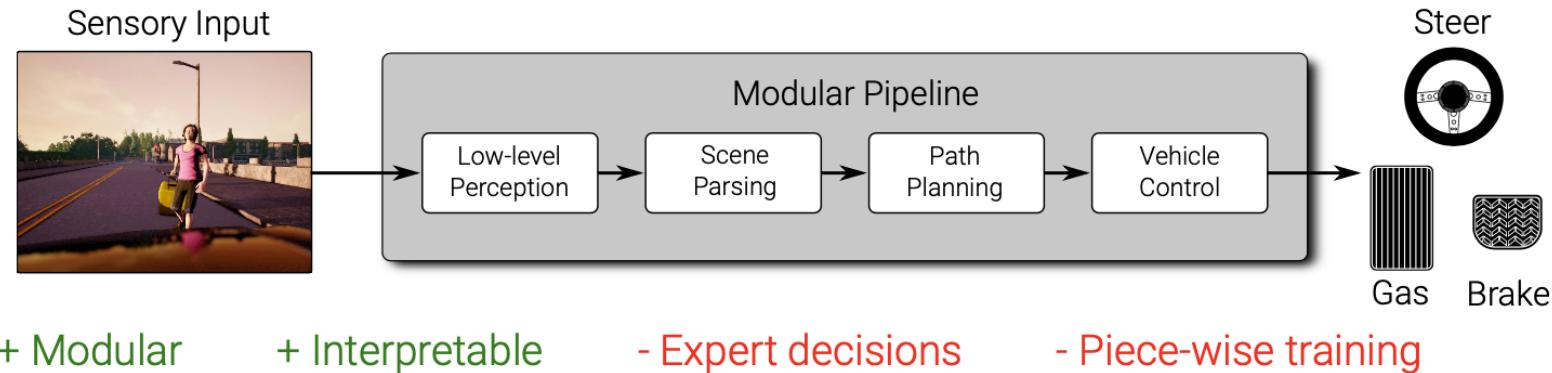
→ Direct perception

- Affordance Learning
- Conditional Affordance Learning
- Visual abstractions



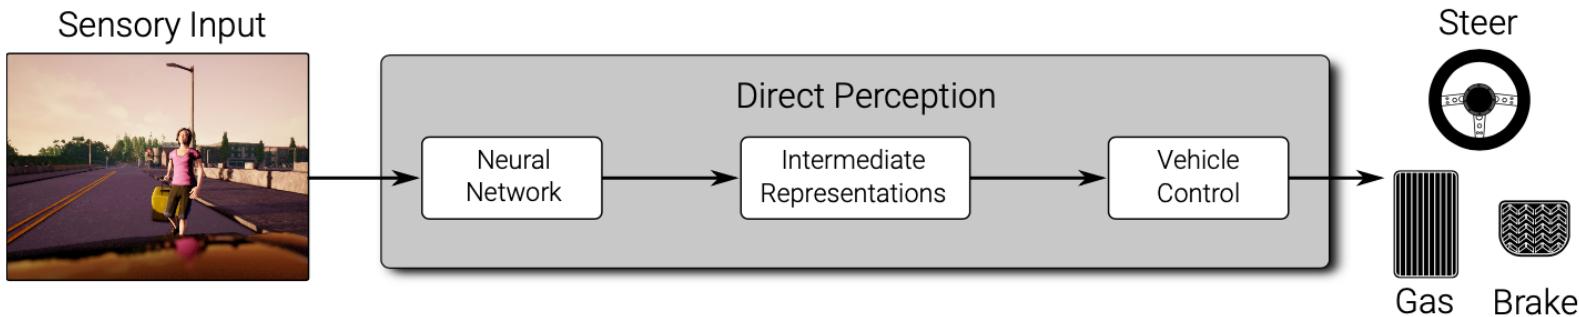
Direct perception

Advantages and drawbacks wrt other methods



Direct perception

A middle ground between end-to-end and modular pipeline

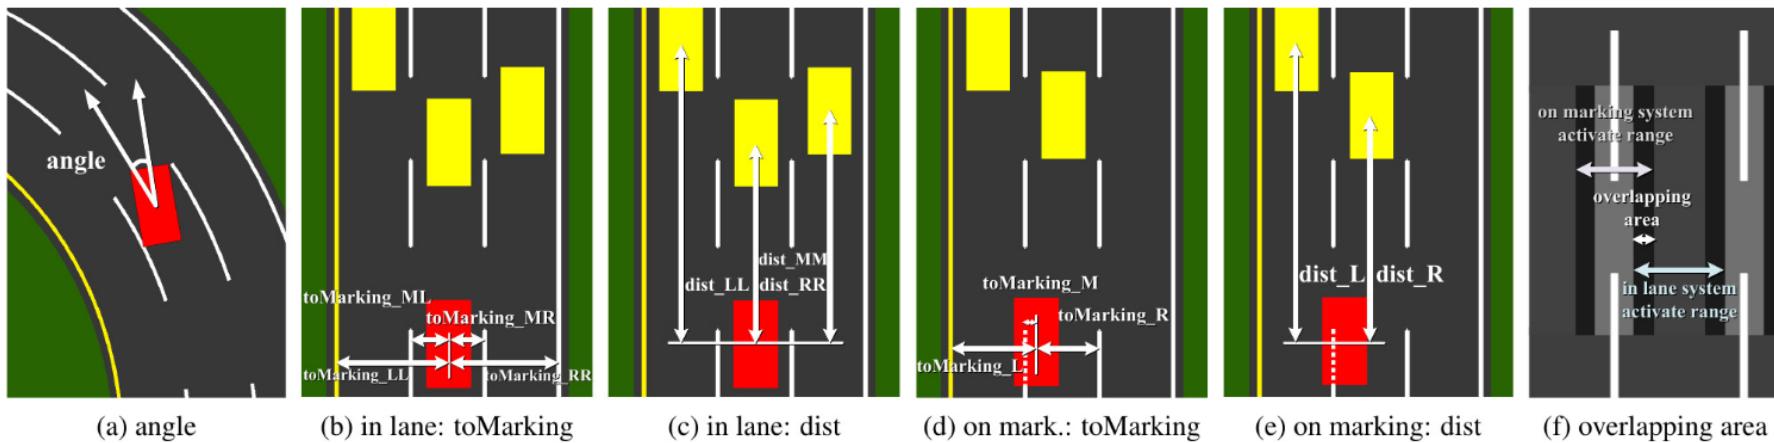


- Hybrid model between end-to-end and modular pipelines
- Learning to predict low-dimensional intermediate representations
- Decoupling perception from planning and control
- Allows to exploit classical controllers or learned controllers (or a mix)

Direct perception for autonomous driving

Using “affordances”

→ Affordance: attribute of the environment which limits the space of actions

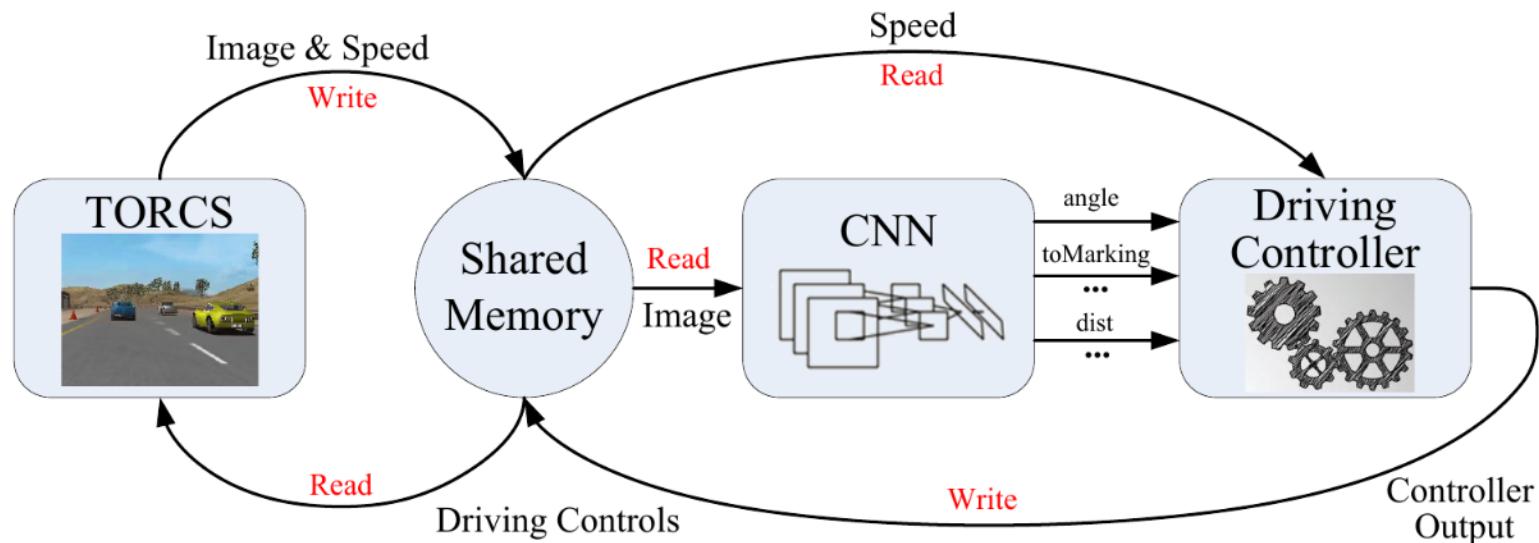


→ In this example, 13 affordances in total

Learning affordances via a CNN

And feeding them to a driving controller

- Using the TORCS simulator to gather data and simulate
- Proposal of a simple CNN (AlexNet)
- Training the affordance indicators
- Using a state machine to drive the car



Controller logic

A state machine for driving the car (using the affordances)

while (in autonomous driving mode)

ConvNet outputs affordance indicators

check availability of both the left and right lanes

if (approaching the preceding car in the same lane)

if (left lane exists **and** available **and** lane changing allowable)

 left lane changing decision made

else if (right lane exists **and** available **and** lane changing allowable)

 right lane changing decision made

else

 slow down decision made

if (normal driving)

center_line= center line of current lane

else if (left/right lane changing)

center_line= center line of objective lane

compute steering command

compute *desired_speed*

compute acceleration/brake command based on *desired_speed*

always:

1) angle: angle between the car's heading and the tangent of the road
"in lane system", when driving in the lane:

- 2) toMarking_LL: distance to the left lane marking of the left lane
- 3) toMarking_ML: distance to the left lane marking of the current lane
- 4) toMarking_MR: distance to the right lane marking of the current lane
- 5) toMarking_RR: distance to the right lane marking of the right lane
- 6) dist_LL: distance to the preceding car in the left lane
- 7) dist_MM: distance to the preceding car in the current lane
- 8) dist_RR: distance to the preceding car in the right lane

"on marking system", when driving on the lane marking:

- 9) toMarking_L: distance to the left lane marking
 - 10) toMarking_M: distance to the central lane marking
 - 11) toMarking_R: distance to the right lane marking
 - 12) dist_L: distance to the preceding car in the left lane
 - 13) dist_R: distance to the preceding car in the right lane
-

Affordance Learning

Results



PRINCETON
UNIVERSITY

Learning Affordance for Direct Perception in Autonomous Driving

Chenyi Chen Ari Seff Alain Kornhauser Jianxiong Xiao

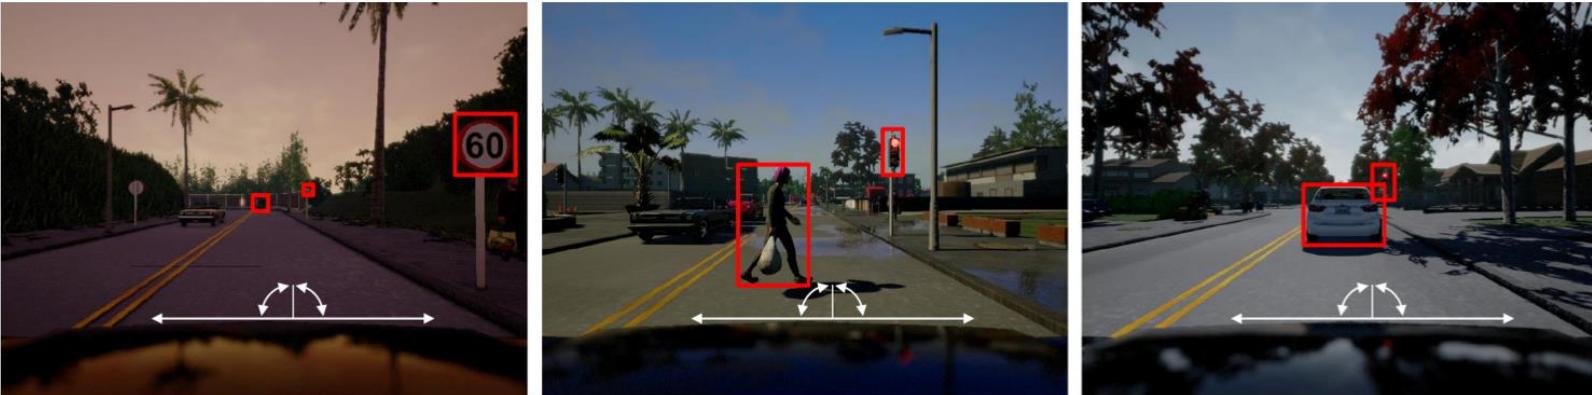
Princeton University



Conditional Affordance Learning (CAL)

Same idea, ported to urban environments and adding a condition

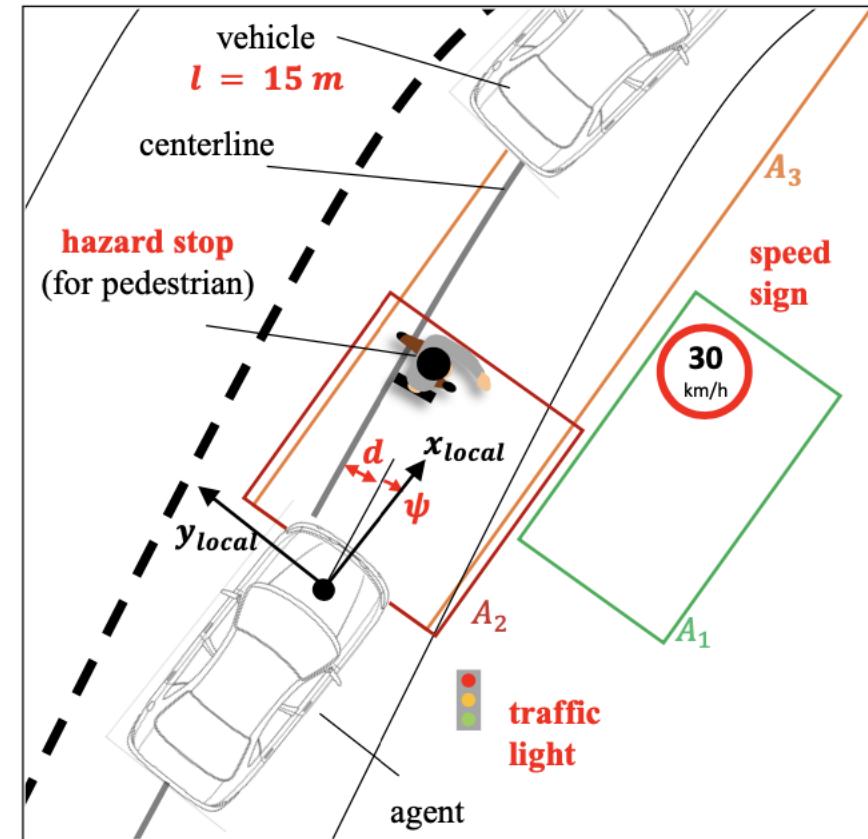
- Using Carla simulator
- Goal: driving from A to B fast, safely and comfortably
- Infractions:
 - driving on wrong lane, driving on sidewalk, running on a red light, violating speed limits, colliding with vehicles, hitting objects/pedestrians



CAL in urban environments

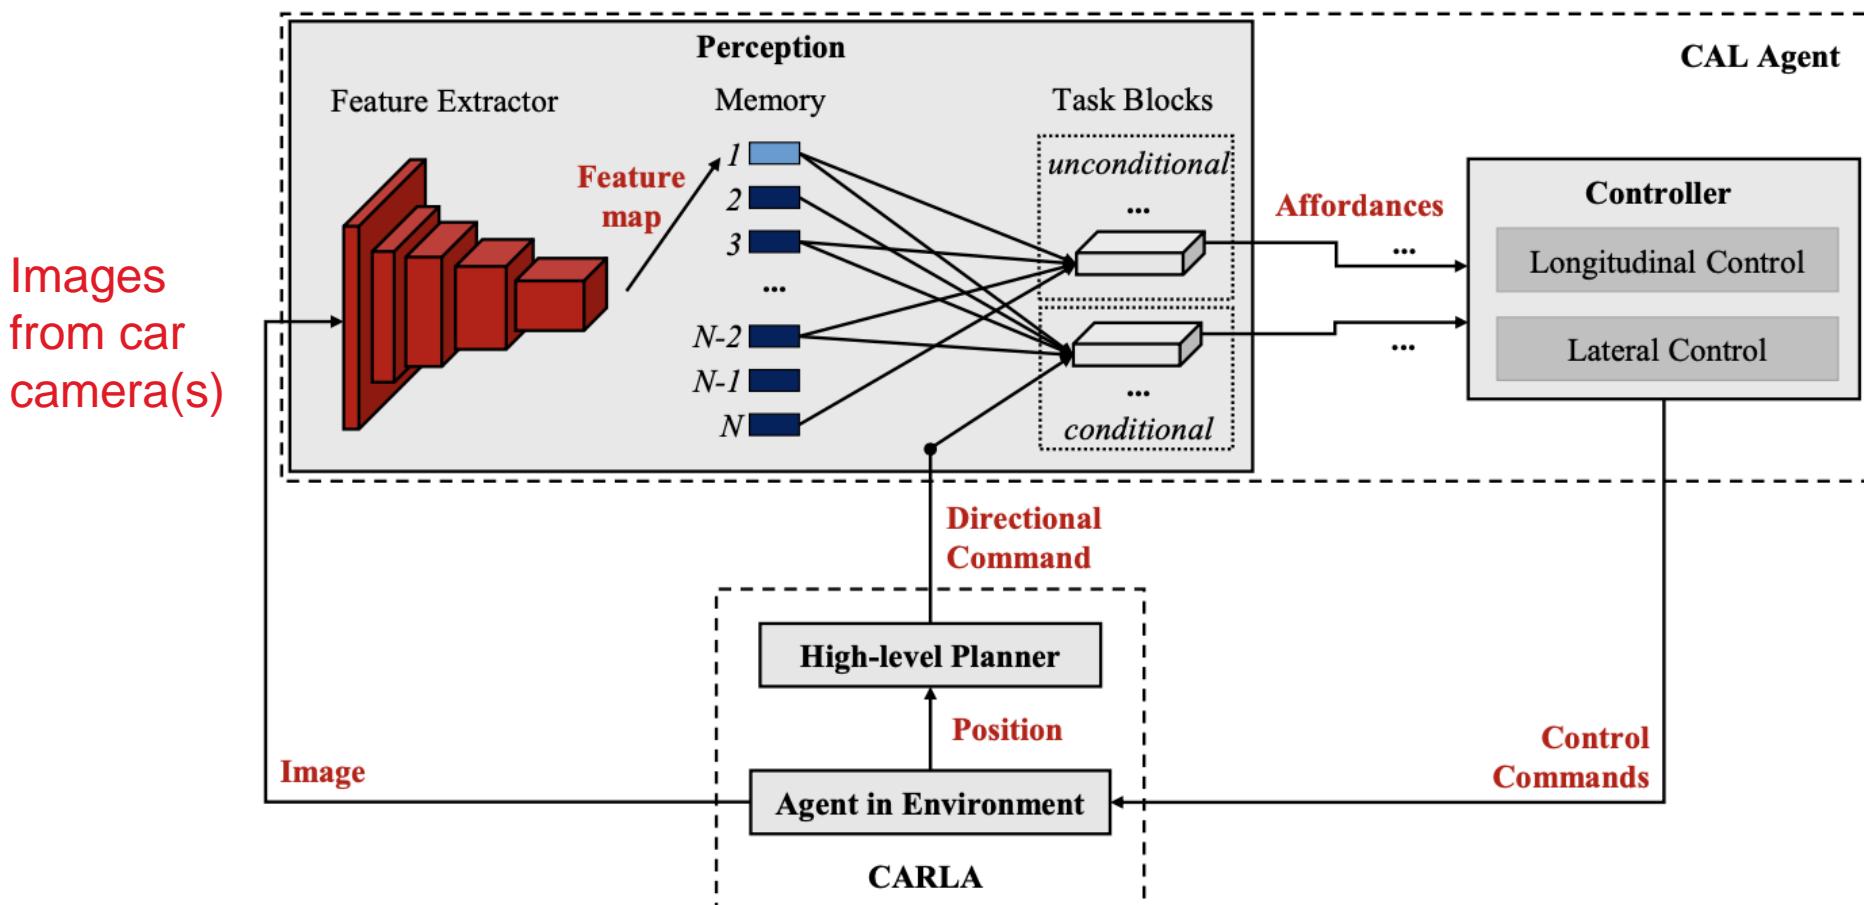
Affordances

- Distance to centerline
- Relative angle to the road
- Distance to lead vehicle
- Speed signs
- Traffic lights
- Hazard stop (pedestrian)



CAL in urban environments

Overview of the solution

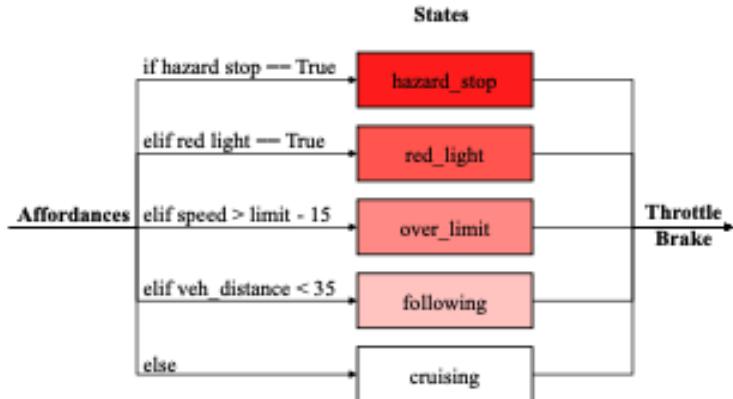


How is the controller block implemented?

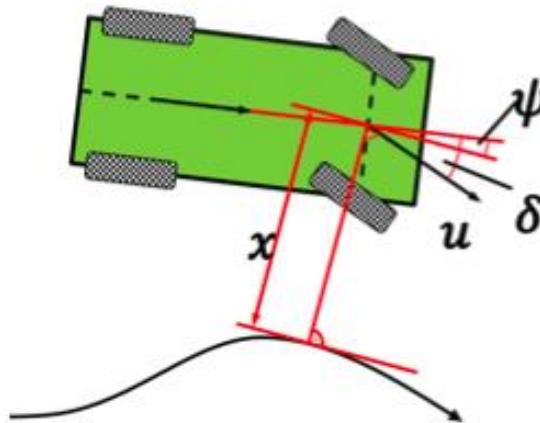
Longitudinal and lateral control

- Can be implemented in “the traditional way”
(vehicle control and dynamics lecture!)
- Longitudinal control
 - Finite State Machine
 - PID controller for cruising
 - Car following model
- Lateral control
 - Stanley controller

Longitudinal Control



Lateral Control



Conditional Affordance Learning

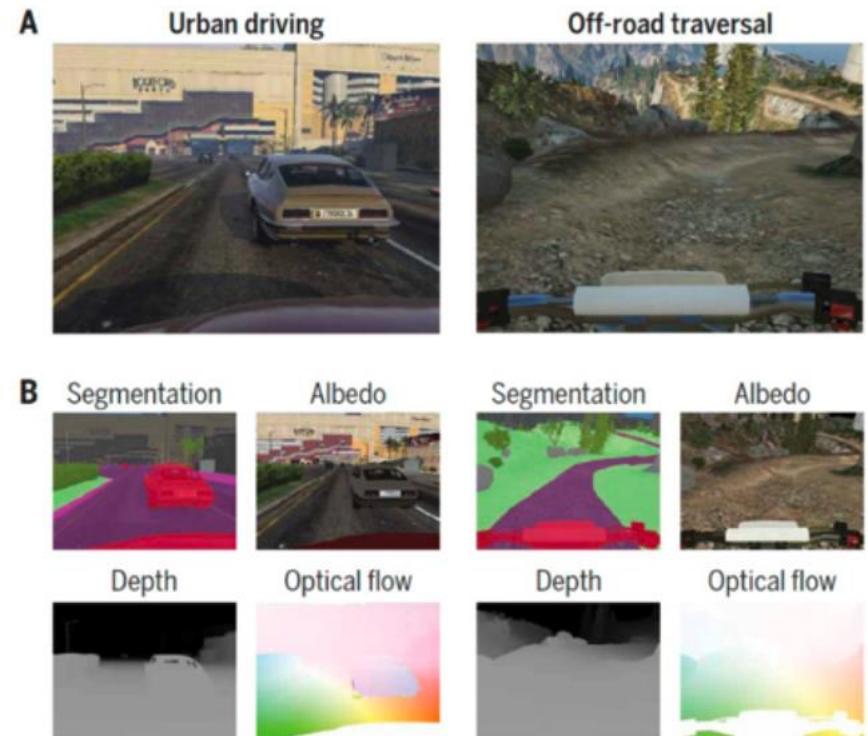
Results



Visual abstractions

Better “intermediate representations” improve results

- Does computer vision matter for action?
 - Yes, it does. Better intermediate representations should improve the results.
- Several different intermediate representations:
 - Image segmentation, depth, optical flow...
- Depth and semantic segmentation often provide better results
 - That is, they generalize better



What is a good visual abstraction?

What does better generalization mean?

- A good visual abstraction is:
 - Invariant : hides irrelevant variations
 - Universal : applicable to a wide range of scenarios
 - Data efficient in memory/computation
 - Label efficient: requires little manual labelling effort
- Semantic segmentation:
 - Encodes task-relevant knowledge and groupings
 - Can be processed with standard CNNs
 - But... labelling time can become huge: ~90min for 1 cityscapes image
- Trade-off when labelling

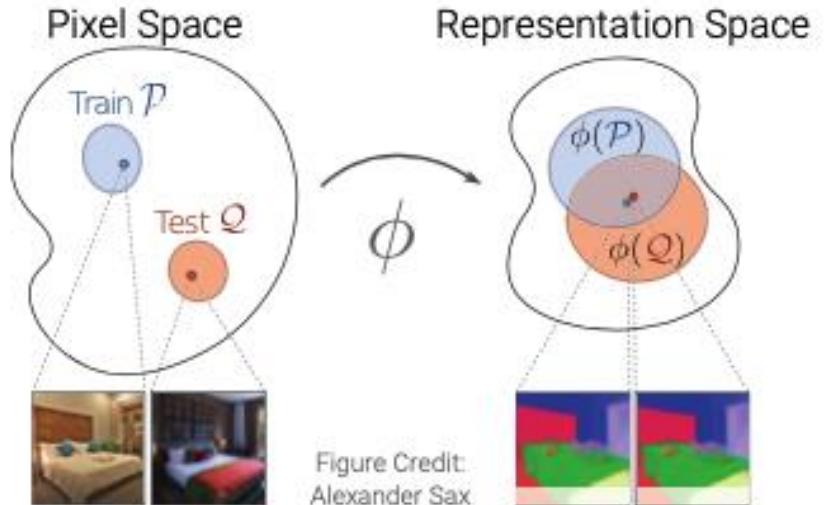


Figure Credit:
Alexander Sax

Identifying the most relevant classes

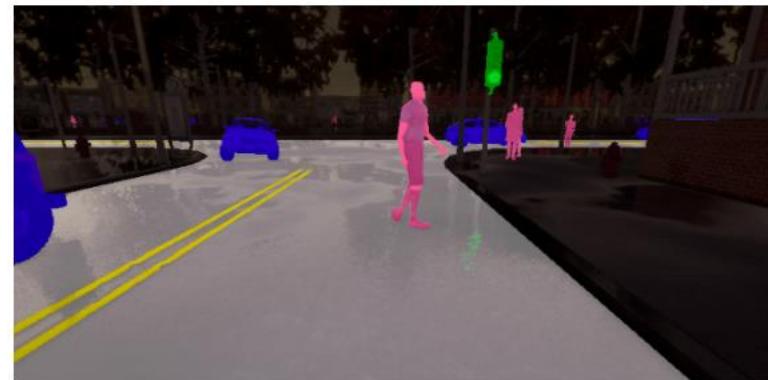
Different number of classes for object detection and segmentation

→ Visual abstraction is an abstraction:

- Algorithm performance not necessarily improves with number of classes



14 classes



6 classes

14 classes: road, lane marking, vehicle, pedestrian, green light, red light, sidewalk, building, fence, pole, vegetation, wall, traffic sign, other.

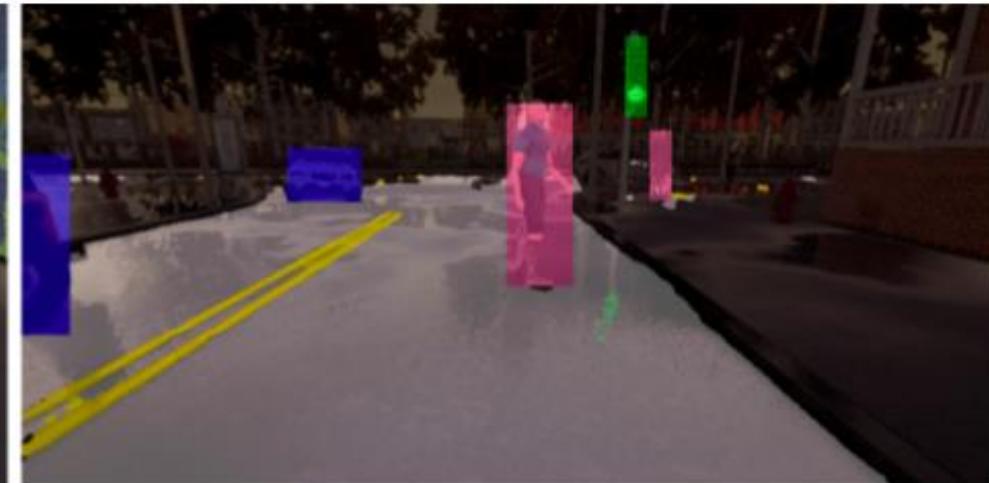
6 classes: road, lane marking, vehicle, pedestrian, greenlight, redlight.

Training time and annotation time

Similar (and even better) success rate for simplified classification



Trained with 6400 finely annotated images and 14 classes
Annotation time ≈ 7500 hours, policy success rate = 50%



Trained with 1600 coarsely annotated images and 6 classes
Annotation time ≈ 50 hours, policy success rate = 58%

TODOs for today

Exercises

1. Analysis of the Dronet running on the Crazyflie
(papers provided on Cyberlearn)
2. Conditional Imitation Learning
3. Conditional Affordance Learning

