

The image shows a modern, multi-story building with a courtyard. The building features large windows with orange frames and shutters. The courtyard is paved with light-colored bricks and has several concrete benches. A few people are visible in the courtyard. The sky is blue with some clouds. The logo "HEIG<sup>VD</sup>" is overlaid in the center of the image.

HEIG<sup>VD</sup>

# Intelligence Artificielle pour les systèmes autonomes (IAA)

**Modular pipeline: Odometry, slam and localization**

Prof. Yann Thoma - Prof. Marina Zapater

*Février 2024*

*Basé sur le cours du Prof. A. Geiger*





# Summary

## Today's lesson

→ **Odometry**

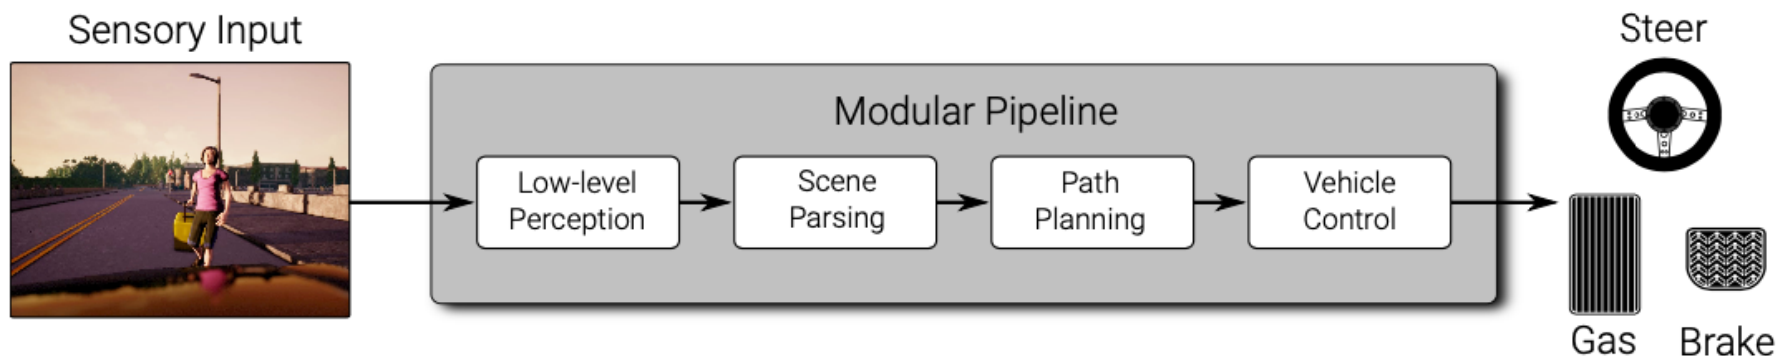
→ SLAM

→ Global localization



# Modular Pipeline

## Reminder of main blocks




- Vehicle control (last week's lecture) (chapter 05)
- **Low-level perception** (today + next lecture)
  - Odometry, motion estimation and localization (SLAM) and global localization methods
- Scene Parsing
- Path planning


# Odometry, motion estimation and localization

## Estimation of ego-motion (own motion) and motion of other traffic participants


- Visual odometry: estimate relative own motion from images
- SLAM (**S**imultaneous **L**ocalization **A**nd **M**apping) algorithms: building a map and simultaneously localize ourselves on that map
- Localization methods: finding global positioning of a vehicle given a map



Visual Odometry



SLAM



Localization

# Odometry

## Using sensors to estimate changes in own motion over time

- Odometry provides relative motion estimates
  - No global position wrt a map
- Sensitive to error accumulation over time
- Can be obtained using a wide variety of sensors
  - **Wheel odometry**: wheel encoders to measure wheel rotation
  - **Inertial Measurement Units (IMU)** measure a body's forces (acceleration)
  - **Visual odometry** : use camera images

A combination of sensors is most often used



Wheel Odometry



Inertial Measurement Unit (IMU)





Visual Odometry

# Wheel odometry

**“The easiest” and most often used in robots**

- Using wheel sensors to update position
  - Incremental encoders : as wheels rotate, output of photodetectors oscillates between low and high. We count the number of pulses
- Using trigonometry to convert number of pulses counted to actual wheel displacement
- Error sources:
  - Limited resolution and pulse-counting errors
  - Unequal (or inaccurate measurement) of wheel diameter
  - Variation of the contact point of the wheel
  - Unequal floor contact, variable friction and slipping






## And in a modern car?

### Speed meter of a modern car

- Speed sensor connected to the car's transmission
- Either mechanical or ferromagnetic
- After the gear box







# Inertial Measurement Units (IMUs)

## Measuring gravity and angular rate

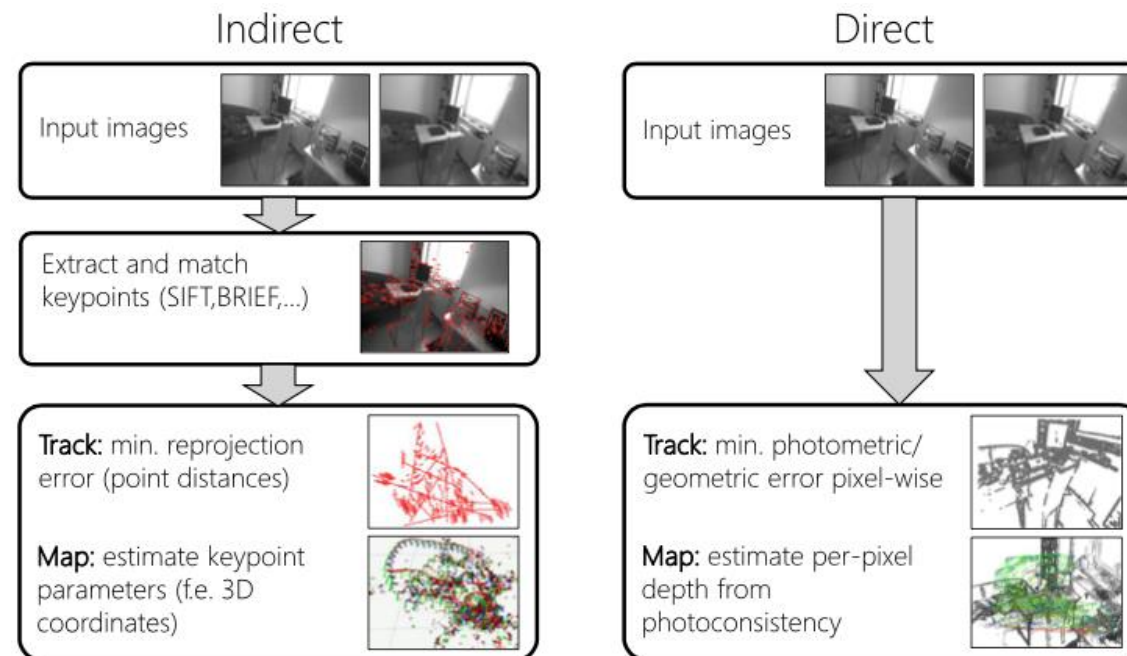
- Gyroscopes: provide a measure of angular rate
  - Twisting or rotational movement
  - Yaw, pitch and roll
  - Can be used to detect object orientation in space
  
- Accelerometers: measure forces/acceleration
  - Measures linear acceleration (changes of velocity) in a single axis
  - Normally IMUs are equipped with a 3-axis accelerometer
  
- Magnetometer (optional)
  - To know the earth's magnetic North
  - Determines heading



# Visual odometry (VO)

## Tracking position and orientation



- Tracks pose (position **and** orientation) of the camera wrt the environment from images
- Considers a limited set of recent images: for real-time and to reduce error accumulation
- A sparse local map is built as a by-product (but is not the main goal)
- Either “Indirect” or “Direct” methods can be used



# Indirect visual odometry

## Extract and match keypoints

- Detect salient points in the images : edges and blobs
  - A blob is a region in which some properties are similar
- Match features between two images by their similarity (find correspondences)
- Requires:
  - Taking perspective into consideration: mapping 2D points into 3D space






# Indirect visual odometry

## Extract and match keypoints

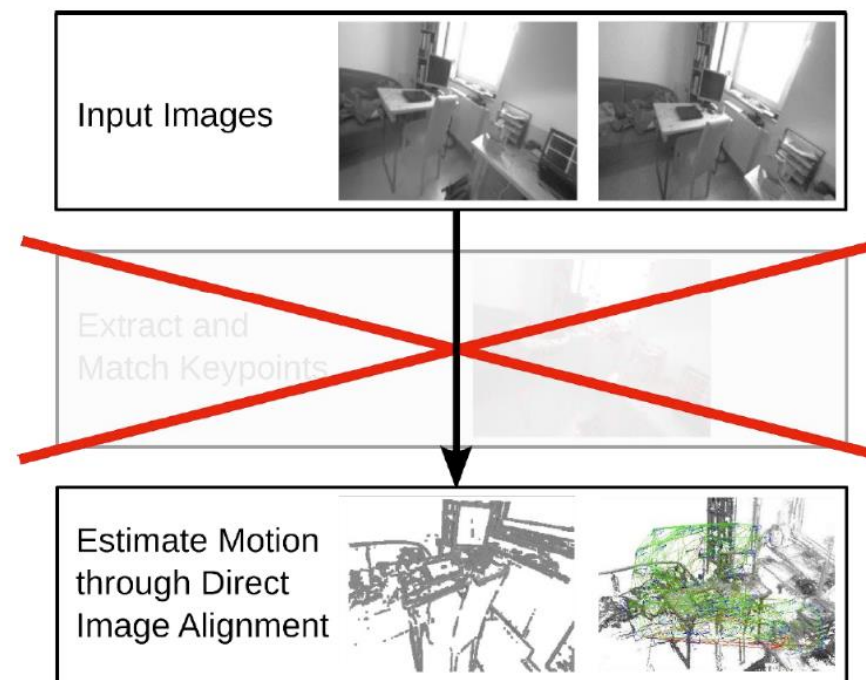

- Detect salient points in the images : edges and blobs
  - A blob is a region in which some properties are similar
- Match features between two images by their similarity (find correspondences)
- Requires:
  - Taking perspective into consideration: mapping 2D points into 3D space



# Direct visual odometry

## Avoiding keypoint detection and matching

- Direct image alignment based on pixels
  - If we know per-pixel depth, we can simulate an image from a different viewpoint
- Depth from sensor (lidar, for example) or through multi-view stereo cameras



# Summary

## Today's lesson

→ Odometry

→ **SLAM**

→ Global localization





# Simultaneous localization and mapping (SLAM)

**SLAM: joint optimization of poses and map**

- Until now: considering of 2 adjacent frames, no focus on map
- Now: considering larger windows → we want poses and location on map
- SLAM is chicken-and-egg problem: localization requires mapping and vice-versa
- SLAM aims to correct accumulation errors via loop-closure detection
- The resulting map is used for localization

# Two different formulations for SLAM

## Full SLAM vs. online SLAM

→ Full SLAM:


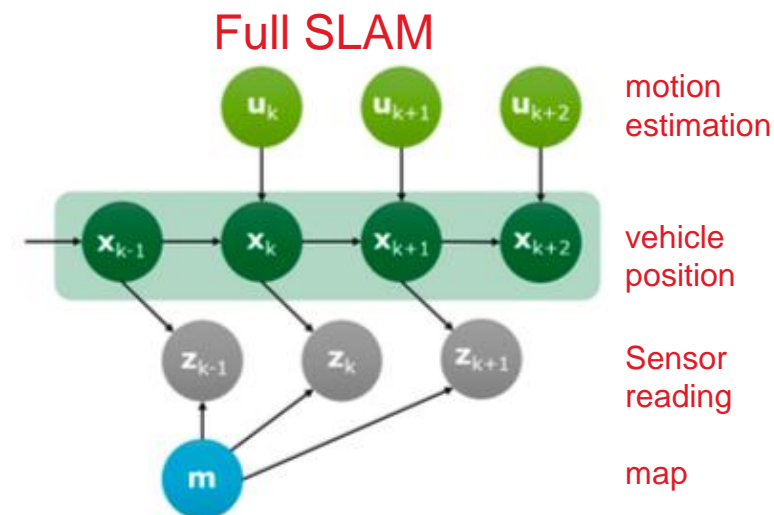
- Estimating the whole trajectory of vehicle and map given all the inputs and measurements
- Very complex! Problem grows very fast

→ Online SLAM (most common today):

- Estimating current position only based on the last sensor information (incremental way)

→ Two common approaches (roughly speaking):

- Filter-based approaches : for online SLAM
- Optimization methods : traditionally used for full SLAM, now used for online SLAM too



[https://hal.science/hal-01615897/file/2017-simultaneous localization and mapping a survey of current trends in autonomous driving.pdf](https://hal.science/hal-01615897/file/2017-simultaneous%20localization%20and%20mapping%20a%20survey%20of%20current%20trends%20in%20autonomous%20driving.pdf)

# Common approaches for online SLAM

## Filter-based and optimization-based methods

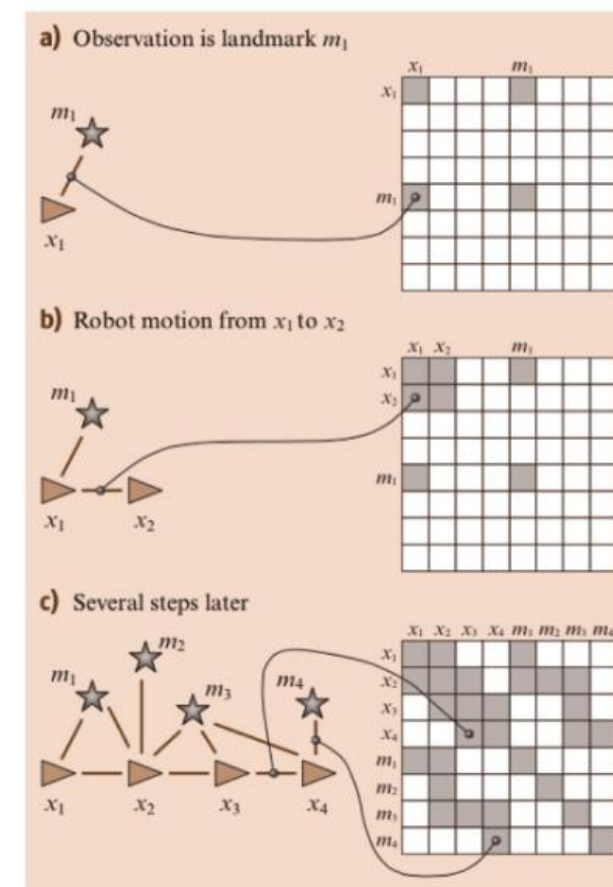
### → Filter-based methods

- Can be used on LIDAR, radar but also in 3D points (vision sensors)
- Example : Extended Kalman Filters - using derivatives of KFs to process data coming from sensors

### → Optimization-based

- Bundle adjustment : jointly optimize 3D structure and the camera parameters
- Graph SLAM: trying to graphically represent the online SLAM problem, and then solve the optimization problem of traversing a graph

### Graph SLAM






# An example of SLAM

## Feature-based SLAM via bundle adjustment

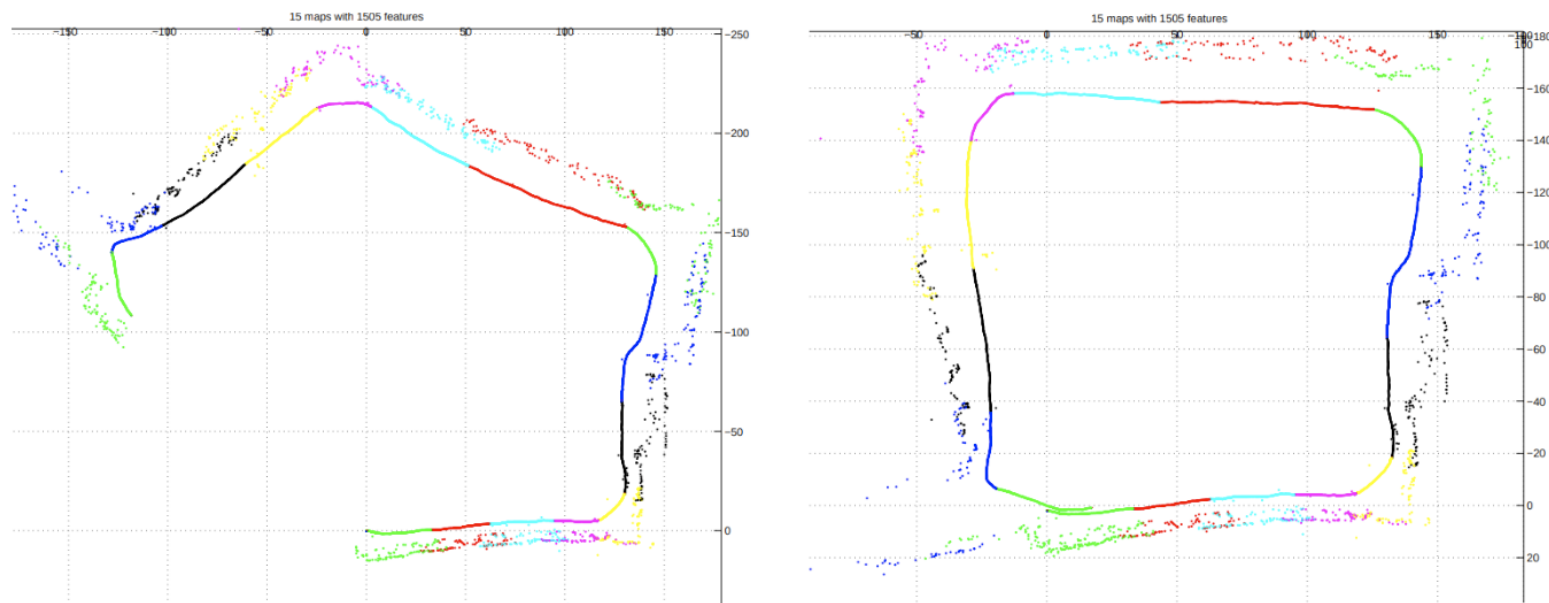
- Optimizing reprojection errors (distance between observed feature and projected 3D point in image plane) wrt camera parameters and 3D point cloud.



# Loop closure detection

Finding correspondences between current and all previous frames

→ And using these correspondences as constraints for optimization



# Many other algorithms for SLAM

A non-exhaustive table for Indirect SLAM methods

MonoSLAM [1]	PTAM [2]	ORB-SLAM 2 [3]
+ monocular cameras	+ monocular cameras	+ monocular cameras + stereo cameras + RGB-D cameras
- no global consistency	- no global consistency	+ global consistency
Extended Kalman filtering of camera pose and point coordinates. Includes a motion model	Mapping as BA over keyframes, real-time tracking towards keyframe	Mapping as BA over keyframes, real-time tracking towards keyframe, loop-closing via place recognition
Filtering	Bundle adjustment	Bundle adjustment
- local accuracy -- global accuracy	+ local accuracy - global accuracy	+ local accuracy ++ global accuracy



# Many other algorithms for SLAM

A non-exhaustive table for Direct SLAM methods

DVO-SLAM [4]	LSD-SLAM [5]	DSO [6]
+ RGB-D cameras	+ monocular cameras + stereo cameras	+ monocular cameras + stereo cameras
+ global consistency	+ global consistency	- no global consistency
camera pose tracking towards keyframe	camera pose tracking towards keyframe	camera pose tracking towards keyframe, camera pose optimization in local keyframe window
+ depth from sensor	+ depth from stereo comparisons & filtering	++ depth optimization in local keyframe window
tracking-only & pose graph optimization	tracking-and-mapping & pose graph optimization	tracking-and-mapping & direct sparse bundle adjustment in local keyframe window with marginalization
+ local accuracy	+ local accuracy	++ local accuracy

# Summary

## Today's lesson

→ Odometry

→ SLAM

→ **Global localization**



# Vehicle Location

## Motivation:

- Control wrt. a path requires knowledge about the position of the vehicle
- Sometimes local knowledge is sufficient (lateral position on highway)
- But often the global location is required (path planning, determining relevant elements that have been marked in a map such as street signs or lane markings)

## Localization Approaches:

- Satellite Localization (uses **infrastructure**)
- Visual Localization (uses **visual map**)
- Map-based Localization (uses **road map**)

# Satellite Localization

## Satellite Systems

- Galileo: 30 satellites (Europe)
- GPS: 24 satellites (United States)
- GLONASS: 24 satellites (Russia)
- BeiDou-2: 30 satellites (China)



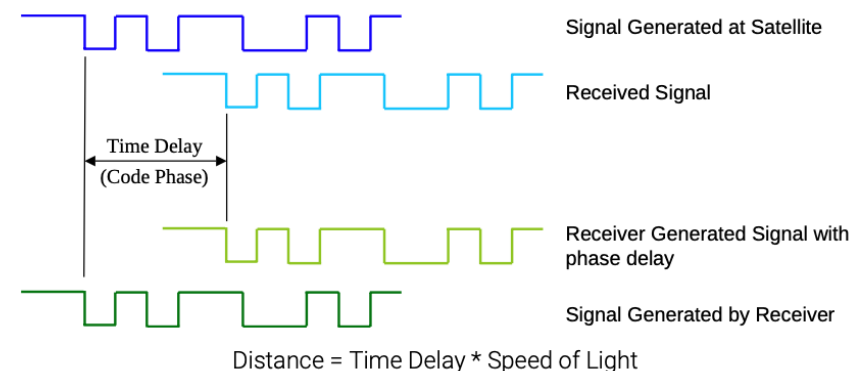
GPS: 4-5 satellites per orbital plane




# Satellite Localization

## Trilateration

- Determine location by measuring distance to satellites at known locations
- Distance to satellite measured as time delay between sent signal/code and received signal/code (receiver knows the code)
- Satellites equipped with an onboard atomic clock
  - 1 satellite: we know the sphere
  - 2 satellites: we know the circle
  - 3 satellites: we know the location (2 possibilities)
  - 4 satellites: time synchronization and exact location
- Differential GPS (DGPS):
  - Improve measurement using ground stations



**GPS equations:** 4 time measurements and 4 unknowns:  $x, y, z, \Delta t$  (see footnote)



# Satellite Localization

## Problems with Satellite Localization

### → Availability

- Satellites not visible in tunnels, narrow city streets
- Dependency on national organizations / interests

### → Accuracy

- 5m-15m for GPS and 0.5m-5m for DGPS
- Only location, no rotation (not full pose)

### → Frequency: max 5-10Hz

### → Sensitive to atmospheric variations and multipath effects (signal reflections)

# Visual localization

## Main idea

- Record map/database of known locations with associated features
  - Features can be extracted from images or laser scans
- Mapping should be conducted under similar conditions as during localization (minimize shift wrt. sensor setup and environmental conditions)
- At localization time, try to retrieve features extracted from the current input image/scan in the map/database
- Either localize only at the image level or refine using triangulation or geometric pose estimation

# Visual localization

Challenges: same image, but appearance and geometry change...



Geometry changes



Appearance changes





# Visual localization

## Topometric localization

- Setup: Record sequence with cameras, Lidar and GPS as ground truth
- Goal: Localize new image (different day/season) wrt. recorded sequence
- Creating a map via a direct graph, where nodes (=frames) are created based on distance threshold.




# Visual localization

## Learning-based localization

PoseNet:

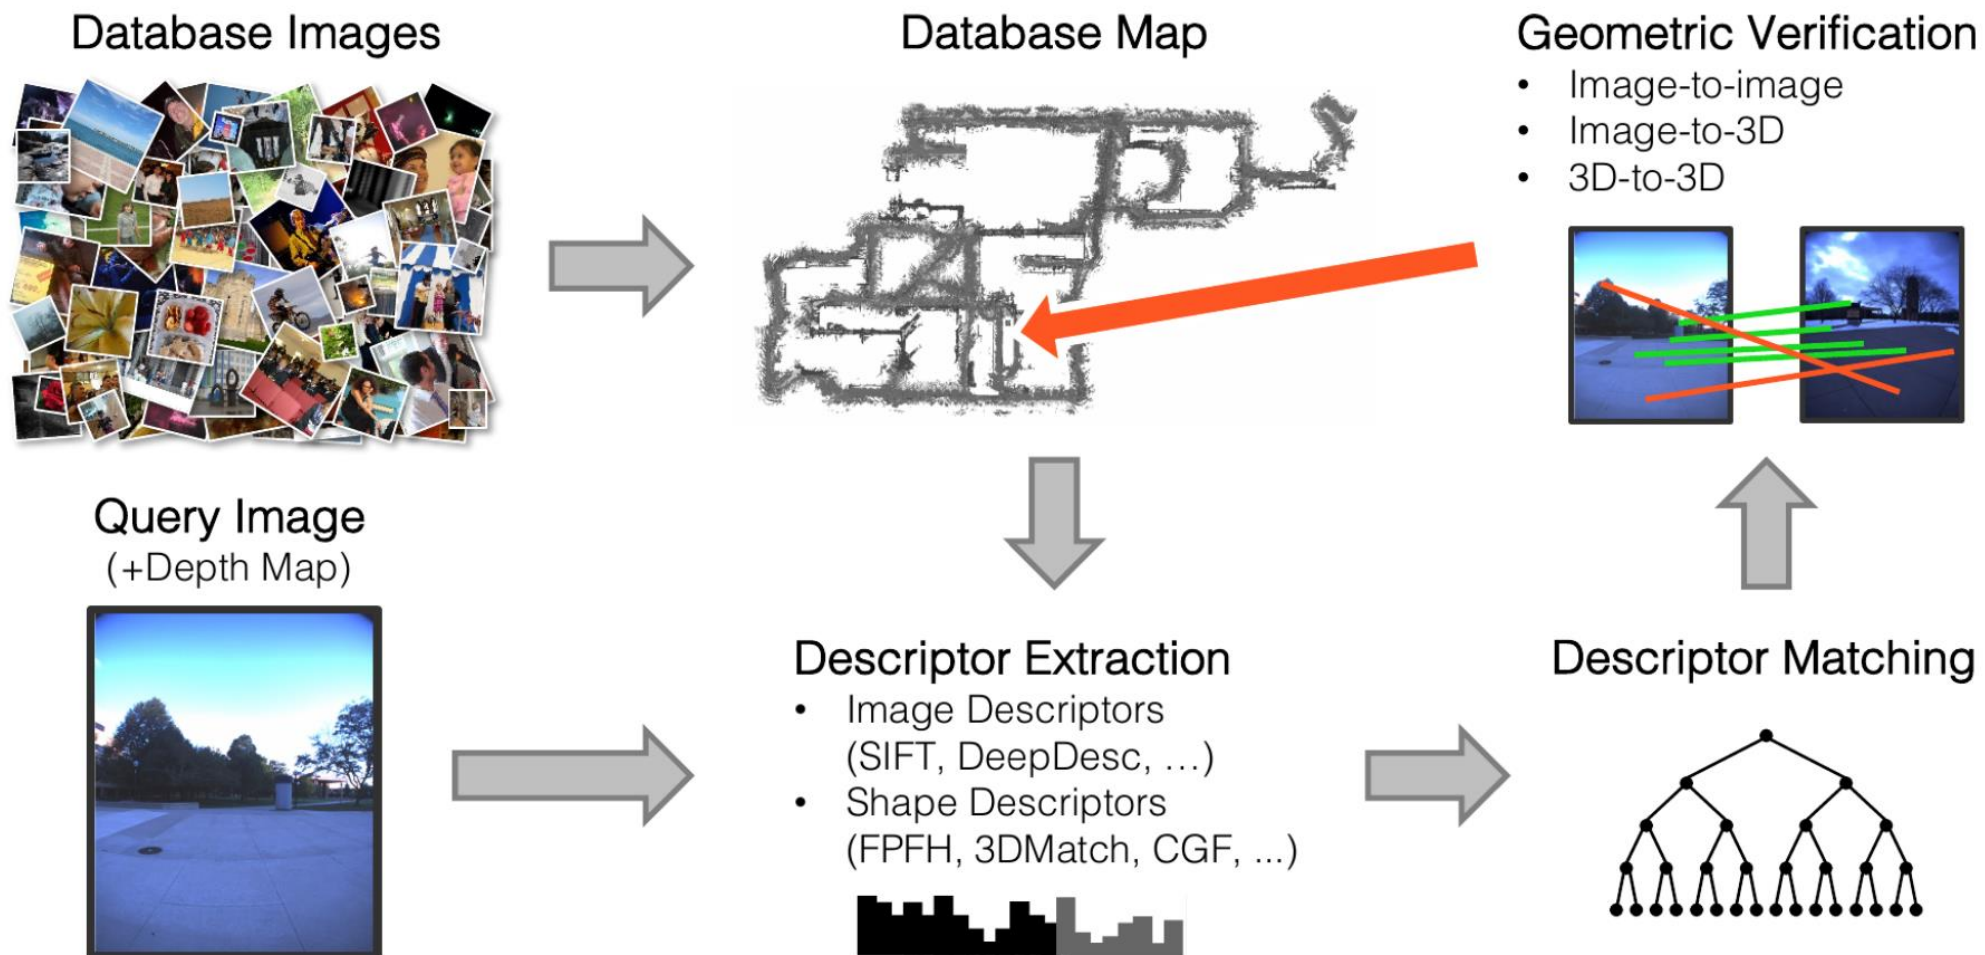
- Input: Single RGB image
- Output: 6 DOF camera pose
- 23 layer deep convolutional network based on GoogleNet
- More robust than feature-based methods
- Less accurate than feature-based methods



(green = train, red = pred, blue = GT)

# Visual localization

## Feature-based localization - Overview

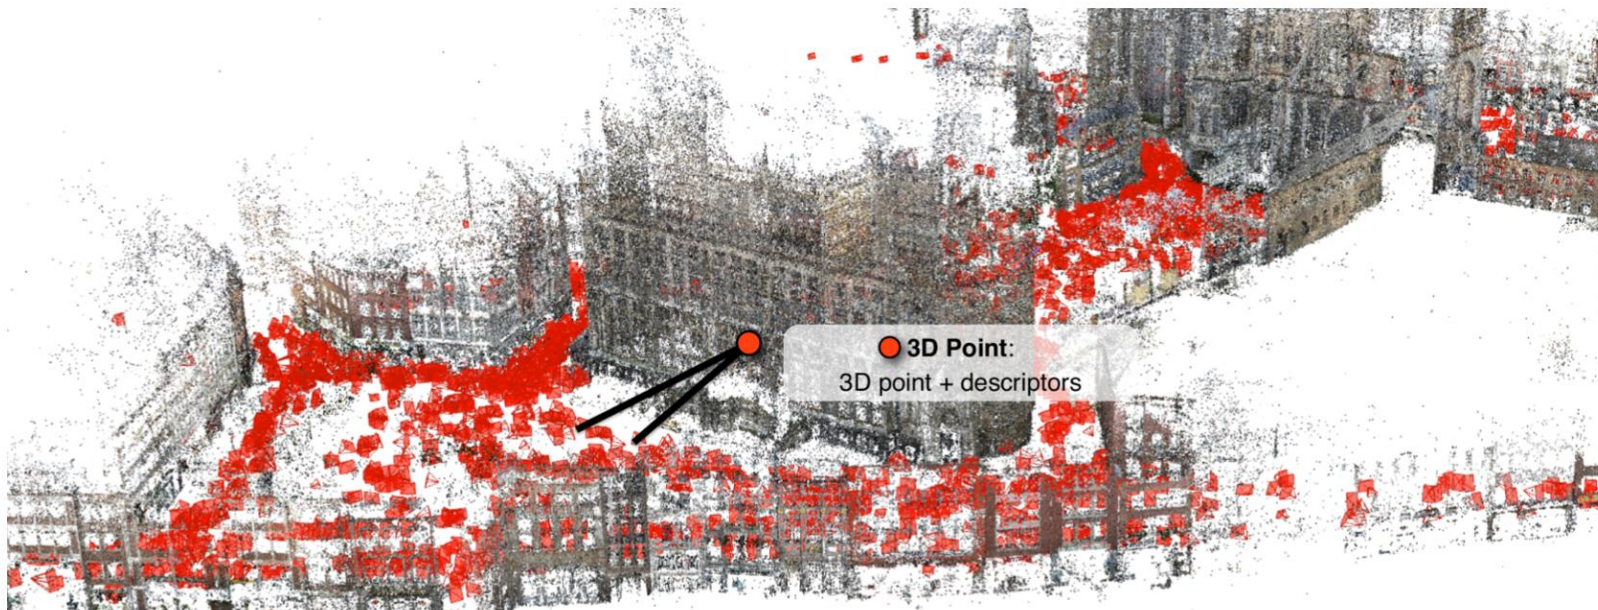




# Feature-based localization

## Mapping

- Sparse 3D reconstruction based on sparse feature correspondences (SLAM)
- Associate each 3D point with a local image descriptor (e.g., SIFT, SURF)






# Feature-based localization

## Mapping


→ Semantic viewpoint invariant feature representations can facilitate matching



# Feature-based localization

## Descriptor matching


- Search nearest neighbours of features from query image (in descriptor space)
- Approximate search due to curse of dimensionality
- Popular approaches: k-dtrees, invertedindex,...
- Good software packages available: FLANN,FAISS,...



# Map-based localization

## Using Visual Odometry (VO)

- Update location probability distribution on map using VO measurements



HE<sup>VD</sup>  
IG

**REDS**  
Institut  
Reconfigurable  
and Embedded  
Digital Systems